


The driver of length-gen failure is attention dispersion off the correct source, not variance collapse
(baselines, order-invariant tasks, both positional encodings, all lengths & seeds)

accuracy vs attention-on-source ($r = 0.94$)



accuracy vs variance collapse ($r = 0.47$)

