

Fixed-Spectrum Swaps Separate Attention Capacity from Assignment

Anonymous submission

Abstract

Attention concentration cannot identify where a transformer routes information: two attention rows may share the same weight spectrum while assigning their largest weights to different tokens. We introduce fixed-spectrum swaps, a causal intervention that preserves every permutation-invariant property of an attention row while changing its assignment to task-labeled positions. Across controlled retrieval, pretrained language models, and context-grounded question answering, the intervention separates selective capacity from correct assignment and reveals assignment-responsive, saturated, and selector-mismatch regimes. A local value-aware utility term ranks finite margin effects in controlled retrieval with 0.989 Spearman correlation and 99.5% sign agreement, and utility-selected circuits yield positive held-out effects on natural question answering. Context-length sweeps show where these diagnoses stop being sufficient: in some cases, behavior degrades without corresponding routing changes. Fixed-spectrum swaps give a controlled test of whether attention assignment causally affects a model’s prediction.

1 Background

Our question is whether attention concentration identifies task-correct routing. At fixed concentration, does assigning the same attention spectrum to task-required evidence change the prediction?

This question matters when performance declines with length. A transformer may learn a retrieval rule on short sequences and then lose it as distractors accumulate. We see this pattern in copying, key-value recall, and algorithmic tasks (Kazemnejad et al. 2023; Wang et al. 2024; Vasylenko et al. 2025). As contexts grow, attention often becomes more diffuse, entropy rises, and attention-output variance can collapse (Li, Boduljak, and Zhou 2025). These trends describe scale and concentration, yet they do not tell us where the attention goes.

The non-identifiability is structural. Two attention rows can share their sorted weights, entropy, maximum, and every norm while placing their largest weights on different tokens. One can favor the evidence and the other a distractor.

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

We test this with a fixed-spectrum intervention. At the answer query, we exchange the weight on a known source with a distractor weight and then run the unchanged model. Because the exchange preserves the complete spectrum, a distractor-only exchange gives a matched control. This paired design also lets us ask when moving weight to the source helps. The answer depends on the spectrum’s selective capacity, its assignment to task-labeled positions, and the downstream utility of the moved values.

Source-max raises per-token accuracy by 0.153, source-min lowers it by 0.043, and the distractor control changes it by 0.011. As the spectrum sharpens, the gap between the two assignments grows even though sharpening the original assignment remains near zero. A local utility term then orders the finite margin effect with 0.989 Spearman correlation and 99.5% sign agreement. Qwen and Pythia show the same direction, while Gemma and SmoLM2 show saturation and selector mismatch. Utility-selected routing also transfers to context-grounded natural QA.


Our contributions are:

- **Fixed-spectrum identification.** We separate selective capacity from task-conditioned assignment by preserving the complete attention spectrum and altering only its token placement.
- **Utility-conditioned mechanism.** We characterize when source-donor reassignment changes the final output and test the prediction against exact interventions.
- **Transfer and boundaries.** We test the intervention across controlled tasks, pretrained models, multi-evidence computation, and natural QA, and identify saturation and selector-mismatch regimes.

2 Related Work

Attention concentration and what it does not identify. Attention concentration has been connected to optimization, token selection, extrapolation, and sparse long-context computation (Tarzanagh et al. 2023; Wang et al. 2024; Nakanishi 2025; Vasylenko et al. 2025). Other accounts emphasize attention-output scale or normalization (Li, Boduljak, and Zhou 2025; Zsámboki et al. 2025). These statistics often move together as context length changes. They are used as evidence in routing accounts, but their status as identifying measures of task-correct routing is generally untested. Our

Use case: Does routing the right evidence change a model's answer?



We preserve the model, prompt, values, and complete weight spectrum, and measure the paired change in target-answer margin ΔM .

Figure 1: Question-answering use case. For the same context and question, a fixed-spectrum swap moves the maximum attention weight from a distractor (“Paris”) to required evidence (“Tokyo”). The model, values, and weight spectrum remain unchanged; the paired target-answer margin tests whether assignment affects the answer.

intervention tests that identification question while the entire spectrum is held fixed.

Attention intervention and non-identifiability. Attention weights have long been studied as explanations, perturbed to test prediction sensitivity, and shown to be non-identifiable without their value vectors (Jain and Wallace 2019; Wiegrefe and Pinter 2019; Brunner et al. 2020; Kobayashi et al. 2020; Hassid et al. 2022). Causal circuit analysis patches internal activations and measures the resulting output change (Yamakoshi et al. 2023; Conmy et al. 2023). Attribution patching uses a gradient approximation to screen many candidate components efficiently (Syed, Rager, and Conmy 2023; Kramár et al. 2024). Our utility term has a related first-order form. Here we evaluate it along an exact permutation of a model’s own attention row, alongside the finite swap itself.

Closest attention interventions. Attention Knockout blocks selected edges and renormalizes the row to trace factual information flow (Geva et al. 2023). Efficient attention projects a row onto its output-effective subspace and studies when attention supports causal explanation (Naim and Asher 2024). Contribution Weights combine attention mass with projected-value magnitude and alignment, then test token importance through removal (Cunningham and Muca Cirone 2026). These approaches delete, block, renormalize, or rank parts of an attention pathway. Our reassignment leaves the complete attention-weight multiset intact and changes its assignment to token positions. Equation 5 then connects that finite, task-labeled swap to the target margin.

Retrieval-based long-context architectures. Sparse and retrieval-based architectures ask how a model can access distant context efficiently (Vasylenko et al. 2025; Hu et al. 2025). Grouped Cross Attention, for example, learns to retrieve useful chunks into a bounded attention window. We study a different object: a trained attention row with a fixed weight spectrum. We isolate the behavioral effect of assigning its available selectivity to task-required evidence.

3 Methods

3.1 Capacity and Assignment

At an answer query, let $a = (a_1, \dots, a_n)$ be an attention row over n token positions, where $a_j \geq 0$ and $\sum_j a_j = 1$. Let $S \subseteq \{1, \dots, n\}$ contain the $r = |S|$ task-required evidence positions, and let $a_{(1)} \geq \dots \geq a_{(n)}$ denote the same weights in descending order. Define

$$C_r^+ = \sum_{i=1}^r a_{(i)}, \quad C_r^- = \sum_{i=n-r+1}^n a_{(i)}, \quad (1)$$

$$K_r(a) = C_r^+ - C_r^-.$$

C_r^+ and C_r^- are the largest and smallest evidence masses available under permutations of the fixed spectrum, and $K_r(a)$ is the available routing range. An assignment is a permutation π of $\{1, \dots, n\}$ that places weight $a_{\pi(j)}$ at token position j . Its evidence mass and normalized assignment coordinate are

$$m_S(\pi) = \sum_{j \in S} a_{\pi(j)}, \quad (2)$$

$$\rho_S(\pi) = \frac{m_S(\pi) - C_r^-}{K_r(a)} \in [0, 1], \quad K_r(a) > 0.$$

Thus $m_S(\pi) = C_r^- + \rho_S(\pi)K_r(a)$. If $K_r(a) = 0$, every assignment has the same evidence mass.

Capacity specifies how much evidence mass the spectrum can support.
Assignment specifies where that mass is placed.
Utility specifies whether moving it changes the answer.

Proposition 1 (Fixed-spectrum capacity-assignment decomposition). *For every assignment π of a fixed weight multiset,*

$$C_r^- \leq m_S(\pi) \leq C_r^+, \quad (3)$$

and both bounds are attained. If $K_r(a) > 0$, Equation 2 defines a unique $\rho_S(\pi)$. No statistic invariant to token permutation can determine ρ_S .

Proof sketch. An r -position evidence set receives at most the r largest weights and at least the r smallest. Assigning those weights to S attains each endpoint. Normalizing evidence mass within this interval gives ρ_S . Any permutation-invariant statistic is constant across all assignments, including assignments at opposite endpoints. The supplement gives the full exchange argument and the corresponding affine margin range. The decomposition is algebraic for every r . Behavioral validation is positive for $r = 1$ and $r = 2$; the $r = 3$ interval includes zero and the $r = 4$ models miss competence, so efficacy at higher arity is unestablished.

3.2 Utility-Conditioned Output Effects

Capacity and assignment describe attention mass. Prediction also depends on values and downstream computation. Let H be the selected attention heads. In each $h \in H$, let s be the evidence position, d_h the donor position whose weight is exchanged with it, and a_{hj} the original attention weight on position j . Define the transferred mass $\delta_h = a_{hd_h} - a_{hs}$. Let x_{hj} be token j 's residual-stream contribution after head h 's value and output projections. If z is the unpatched residual state, the intervention gives $z' = z + \Delta z$, where

$$\Delta z = \sum_{h \in H} \delta_h (x_{hs} - x_{hd_h}). \quad (4)$$

Let $\ell_k(z)$ be the downstream logit for token k , with target y and competitor c . Define $M_c(z) = \ell_y(z) - \ell_c(z)$, the exact change $\Delta M_c = M_c(z + \Delta z) - M_c(z)$, and local utility $u_{hj}^{(c)} = \nabla_z M_c(z)^\top x_{hj}$.

Proposition 2 (Utility-conditioned swap effect). *If $\nabla_z M_c$ is L_c -Lipschitz on $\{z + t\Delta z : t \in [0, 1]\}$, then*

$$\Delta M_c = \sum_{h \in H} \delta_h (u_{hs}^{(c)} - u_{hd_h}^{(c)}) + R, \quad |R| \leq \frac{L_c}{2} \|\Delta z\|_2^2. \quad (5)$$

For an affine downstream margin, $R = 0$.

The Taylor remainder is standard. The new component is the task-labeled capacity-assignment coordinate coupled to the exact fixed-spectrum swap vector in Equation 4.

The intervention transfers attention mass from a donor token to the source. It helps when the source contributes more positively to the target margin than the displaced donor, unless downstream curvature overwhelms that gain. Source-max can therefore succeed, saturate, or fail under the same attention-mass rule. The supplement derives Equation 5, extends it to evidence sets, and states a sufficient margin certificate.

4 Experimental Setup

Controlled tasks. Argmax retrieval presents key-value pairs and requests the value paired with the largest key. Flag retrieval presents values with exactly one marked position and requests that value. Both have one labeled source. Two-evidence modular sum marks two digits and requests their sum modulo ten, so both sources are required. Four-layer decoders are trained with NoPE and RoPE and four seeds

per cell; evaluation reaches $50\times$ training length. An eight-layer, width-512 grid checks scale. The inverse temperature $\beta \in \{1, 2, 4\}$ changes attention capacity while each assignment comparison preserves the resulting sorted spectrum.

Pretrained and natural-language tasks. In-context key-value recall lists pairs followed by a query key; the paired value is the labeled source. We test Qwen2.5, Pythia, Gemma, and SmolLM2 across increasing pair counts. The selector audit compares circuits ranked on disjoint calibration data by source mass or by output-conditioned utility. Natural MCQA uses four-choice questions with one gold passage and unrelated passages. Questions must be correct with the full context and gold passage alone, but incorrect without context. This filtering defines a model-specific context-dependent subset and is not intended to estimate average performance over the full dataset. Its length ladder freezes the four-passage circuit before forming nested 8-, 16-, and 32-passage contexts with additional unrelated passages.

Interventions and estimation. The distractor control swaps the largest and smallest non-source weights, preserving source mass while permuting the same attention row. For a source position s , let $d^+ = \arg \max_j a_j$, $d^- = \arg \min_j a_j$, and $D = \{j : j \neq s\}$ over valid tokens. Source-max swaps (s, d^+) and source-min swaps (s, d^-) . In controlled and pretrained key-value recall, the distractor control uses the maximum and minimum within D . Natural QA uses a displacement-matched control. With $\delta = a_{d^+} - a_s$, it chooses

$$(p, q) = \arg \min_{\{j, k\} \subset D} |a_j - a_k| - \delta \quad (6)$$

and swaps (p, q) . Among all non-source pairs, this choice minimizes the absolute difference between the control's L_1 displacement and the treatment's L_1 displacement; the remaining difference is the minimum attainable mismatch for that row. It holds source mass fixed, and every control preserves the complete spectrum. For natural QA, displacement is matched first because it measures perturbation magnitude. The control does not match the donor token, positional distance, or value norm, which remain separate confounds. Calibration and evaluation examples are disjoint, circuits are frozen before paired evaluation, and hierarchical intervals resample calibration or model seeds above paired examples. We interpret a cell only after its frozen competence gate passes. Three confirmatory contrasts form one inferential family: controlled source-max versus distractor control, the capacity-by-assignment interaction, and natural utility-selected source-max versus matched control. We use two-sided exact sign-flip tests at the trained-model or calibration-seed cluster level and Holm correction across the family. Other intervals are descriptive and unadjusted. These swaps are counterfactual probes and need not lie on the model's endogenous query-key manifold. Treatment and control share this intervention class, so their contrast measures downstream sensitivity to assignment, rather than the frequency with which the model naturally creates a particular row. Layer controls (+1.107 for selected $K = 8$ versus -0.059 for a random layer) and coverage dose response show

greater sensitivity under this probe. They do not show that normal query-key computation produces swapped rows or that selected heads are the exclusive retrieval mechanism. In untouched distractor-order variants, the conditional source-mass coefficient is $+1.250$ ($p = .03125$), while the closest-spectrum paired effect is $+0.026$ ($[-.017, +.070]$, $p = .25$). This connects the intervention to an endogenous association without establishing endogenous causality.




Figure 2: Paired causal protocol. A circuit selected on disjoint calibration data is frozen before treatment and control are applied to the same evaluation row. Both paths preserve the spectrum; the natural QA control additionally minimizes displacement mismatch in Equation 6.

Measured outcomes. We record task accuracy, target-versus-competitor margin, source mass, attention-output variance, entropy, norms, and the full sorted spectrum. Every fixed-spectrum condition preserves the selected attention rows to numerical precision. The supplement reports prompts, hyperparameters, per-length tables, competence failures, and preregistered decision rules.

5 Results and Discussion

5.1 Assignment Changes Accuracy at Fixed Spectrum

We begin with the simplest test: move the largest available weight to the source, move it away, or exchange two distractors. Across sixteen trained models, the mean changes are $+0.153$, -0.043 , and $+0.011$, respectively (Figure 3). Source-max beats the control in every model-length comparison. Their primary cluster-level contrast is $+0.142$ (95% CI $[+.080, +.206]$; exact $p = 3.05 \times 10^{-5}$; Holm-adjusted $p = 9.16 \times 10^{-5}$), while complete-spectrum invariant error stays below 10^{-6} .

The controlled audit contains 8,192 head-example pairs. Individual swaps are vacuous in 18.5% of them, but only

86 of 1,024 complete eight-head circuits (8.4%) are entirely vacuous. Mean exact margin change is $+3.891$ over all circuits and $+4.248$ (95% CI $[+3.863, +4.640]$) over the 938 active circuits. The median per-head transferred mass is .025 (90th percentile .930), and effect size rises across transfer quartiles. The supplementary displacement audit reports the full distribution.

The effect grows with circuit coverage (Figure 3b). Patching one, two, four, and eight selected heads gives mean source-max gains of 0.007, 0.013, 0.056, and 0.153 and is monotone in 31 of 32 model-length comparisons, consistent with a distributed effect under the probe.

Source attention tracks accuracy more closely than attention-output variance. Normalizing that variance leaves accuracy unchanged, whereas forcing mass onto the known source restores controlled retrieval. Sharpening the original row has little effect.

The ordering survives the eight-layer replication, both positional encodings, four training seeds, and two controlled tasks. Random-layer and adjacent-layer interventions remain near zero when the selected circuit is held fixed. Full model-level intervals and the variance and patching results appear in the supplement.

5.2 Capacity Amplifies Assignment

We next vary spectrum and assignment independently. Temperature scaling changes the maximum available weight, and at every temperature we evaluate identity, source-max, source-min, and distractor control on the same examples. The source-max minus source-min contrast grows from 0.196 to 0.346, while sharpening the identity assignment changes accuracy by at most 0.009 (Figure 4a). Capacity sets the possible leverage, while assignment determines where that leverage lands. The primary interaction is $+0.150$ (95% CI $[+.094, +.210]$; exact $p = 6.10 \times 10^{-5}$; Holm-adjusted $p = 1.22 \times 10^{-4}$).

For held-out examples we differentiate the target margin along the exact source-max path. The first-order term in Proposition 2 predicts the finite margin change with 0.917 Pearson correlation, 0.989 Spearman correlation, and 99.5% sign agreement (Figure 4b). The remaining error measures downstream curvature along the swap.

Modular sum requires the model to route and combine two marked values. At $5 \times$ training length, assigning the top two weights to evidence raises exact match by 0.066, assigning the bottom two lowers it by 0.291, and evidence-max exceeds a control that preserves both the spectrum and total evidence mass by 0.043 (95% CI $[0.017, 0.072]$). The result therefore extends beyond copying one token. A registered three-evidence estimate includes zero, and four-evidence models miss the frozen competence gate. Higher arity remains unresolved.

Evidence-max and evidence-min preserve the complete spectrum. The multi-evidence control also preserves total evidence mass, changing only which evidence token receives each weight. Its small effect separates aggregate source mass from useful evidence assignment.




Figure 3: Assignment changes retrieval despite an identical attention spectrum. Source-max improves accuracy, source-min harms it, and the distractor-only control remains near zero (a). The effect increases with the fraction of the selected circuit receiving the intervention (b). Diamonds and intervals summarize trained models; faint paths show individual cells.




Figure 4: Capacity controls leverage, and utility predicts the finite intervention. Sharper spectra increase the correct-versus-wrong assignment contrast while sharpening alone remains near zero (a). The first-order utility term predicts the exact source-max margin effect across tasks and positional encodings (b). Error bars summarize trained models.

5.3 The Same Intervention Produces Three Pretrained Regimes

The pretrained models show why the utility term matters. Qwen and Pythia gain margin when source-max replaces the distractor control. Gemma stays accurate as context grows and gains little from extra source mass, even though source-min reliably harms it. This is the saturation regime predicted by Proposition 2: removing evidence matters, while extra routing has little left to add. A quantized Qwen2.5-7B run preserves the positive margin direction.

SmolLM2 provides the contrasting case. Its preregistered source-mass selector fails, yet ranking heads on disjoint calibration examples by output-conditioned gradient or utility restores a positive effect. A ten-seed ablation ranks source

gradient, utility gain, and utility gap above source mass and transfer mass. Rankings vary by architecture, but circuit selection must account for the output.

Interpolating from the natural row to the utility-selected swap gives a monotone margin path over five calibration seeds. On context-grounded natural QA, utility-selected routing improves margin over a displacement-matched control, while source-mass selection harms it. Across six disjoint calibration seeds, the utility-selected contrast is $+0.378$ (hierarchical 95% CI $[+0.253, +0.520]$; exact and Holm-adjusted $p = .03125$). All competence gates pass. An unconstrained generation audit finds no first-answer or repetition cost.

The selector results judge a selector by the frozen circuits it finds and their paired margin effects. They do not test whether the per-row utility functional predicts a finite swap




Figure 5: Matched pretrained-model comparison (seed 0; $n = 128$ paired examples per model-length cell). All panels use the 5-, 20-, and 80-pair ladder and report source-max, source-min, and untouched margin relative to the same distractor control. Qwen and Pythia respond positively, Gemma is saturated, and the source-mass-selected SmolLM2 circuit does not benefit. Error bars show paired 95% intervals.

effect. That held-out predictor audit is reported separately in the supplement.

Clean-to-corrupt activation patching rescues $+0.308$ margin but only between 2.5% and 3.3% of median corruption damage. It restores selected-head outputs after source-value corruption, so its magnitude is not directly comparable to reassignment on an unchanged input.

5.4 Natural Context Marks the Scope Boundary

The fixed-context natural-QA result establishes causal routing in ordinary text. The length ladder is a scope test: it asks whether the same fixed-context protocol remains informative as context grows. We select the circuit at four passages, freeze every choice, and append unrelated passages to form nested 4, 8, 16, and 32-passage contexts. The matched figure fixes seed 0 and resamples paired questions; repeated passage observations from one circuit are not treated as independent models. The supplement reports the additional Qwen seed separately.

Both Qwen and SmolLM2 lose margin and accuracy by 32 passages, yet their source-mass trajectories diverge. Qwen source mass rises as intervention rescue weakens. SmolLM2 source mass falls while rescue stays flat. The shared degradation cannot be attributed to one consistent source-mass trajectory. This is a scope limit of the diagnosis, not a failed prediction.

The protocol has a narrower interpretation. At a fixed context, a spectrum-preserving swap measures whether a chosen circuit can use reassigned evidence. As contexts grow, representation quality, distractor semantics, answer competition, and computation outside that circuit can also matter. The

theorem concerns a specified swap at a fixed context; this experiment does not resolve the other long-context bottlenecks. Exact 32-minus-4 changes and hierarchical intervals appear in the supplement.

The arity checks provide another limit. The two-evidence task passes and supports the mass-preserving control; the three-evidence interval includes zero, and the four-evidence model misses competence. Higher arity remains unresolved. An eight-token generation audit also leaves first-answer accuracy and repetition unchanged, so simple format collapse is not driving the result.

5.5 Identification and Replication Audits

Replication checks vary the calibration sample, patch budget, and model family. Across three frozen Qwen calibration seeds, selected circuits are positive at every head budget: $K = 2, 4, 8$ gives $+0.460$, $+1.054$, and $+1.107$ margin. Adjacent layers are smaller and inconsistent ($+0.147$ and -0.131), while random layers average near zero (-0.059). Full intervals and seed standard deviations are reported in the supplement.

A separate replication tests the utility relation. Circuits are chosen by source mass on disjoint calibration examples; utility is then computed on held-out evaluation rows. Across three seeds per family, every competent source-max cell is positive, and predicted effects rank the exact effects within each cell. Mean Spearman correlations are $.866$ for Qwen2.5-1.5B, $.544$ for Pythia-1.4B, and $.554$ for Gemma-2-2B, with sign agreements $.805$, $.807$, and $.714$. The approximation is imperfect at large effects, so we use it as a local ranking signal rather than a calibrated finite-effect estimator. For the 32 controlled cell means in Figure 4b, the actual-on-predicted

Fixed-spectrum assignment		n	Units	Unpatched	Δ source-max	Δ source-min	Δ distractor	max – distractor
Single evidence	$\beta = 1$	8,192	Acc.	.619	+.153	−.043	+.011	+.142
	$\beta = 4$	8,192	Acc.	.619	+.298	−.048	+.025	+.273
Two evidence	1 ×	2,048	EM	.952	+.019	−.559	+.012	+.007
	5 ×	2,048	EM	.484	+.066	−.291	+.023	+.043
Pretrained KV	Qwen2.5-1.5B	384	Margin	−3.048	+1.990	−.961	+.192	+1.799
	Pythia-1.4B	384	Margin	−3.805	+.826	−1.022	+.146	+.680
	Gemma-2-2B	384	Margin	+2.234	+.040	−1.373	+.013	+.027
	SmoLLM2-1.7B	384	Margin	−2.782	−.168	−.112	−.077	−.091
Output-conditioned circuit selection		n	Units	ΔM versus matched control		ΔM source mass	$\Pr(\Delta M > 0)$	utility – source
Task	Model			Utility gain	Utility gap			
Pretrained KV	Pythia-1.4B	384	Margin	+.519	+.359	+.006	.810	+.513
	SmoLLM2-1.7B	384	Margin	+1.253	+.859	+.006	.969	+1.247
	Qwen2.5-1.5B	384	Margin	+1.493	+1.493	+.069	.956	+1.424

Table 1: Benchmark summary. The pretrained fixed-spectrum rows use the matched 5-, 20-, and 80-pair ladder from Figure 5, with 128 paired examples per model-length cell and $n = 384$ per model. Units make upper-block baselines comparable only within task. The lower block reports the three model families with completed matched selector runs; a Gemma selector-family run was not conducted. It has no unpatched baseline and evaluates selector quality using fixed-spectrum control effects, not per-row utility prediction; the supplementary held-out audit evaluates that prediction. $\Pr(\Delta M > 0)$ is the fraction of paired evaluation examples with a positive utility-gain source-max minus matched-control margin. Bold final columns remain identifiable in grayscale. Full cells and 95% intervals are in the supplement.




Figure 6: Matched natural-QA length ladder (seed 0; $n = 128$ paired examples per model-length cell). From 4 to 32 passages, accuracy falls (a), source routing moves in opposite directions (b), and frozen-circuit rescue stays flat or weakens (c). Bands are paired-example bootstrap 95% intervals.

fit is $1.387x - .102$ (Pearson .978), which deviates from identity calibration.

The exact Qwen rerun keeps precision, lengths, contrast, and sample counts fixed. All five new seeds are positive (mean $+.042$); the combined six-seed test gives $p = .03125$, while the original $+1.882$ magnitude is atypical.

Output-conditioned selectors lead in each family, although their internal ranking changes by architecture. Full selector results, competence-gate failures, and the anonymized code package appear in the supplement.

6 Limitations

Fixed-spectrum swaps identify the causal effect of a specified attention reassignment, not a complete account of model

computation. The protocol requires identifiable evidence positions and a chosen intervention circuit; output-conditioned selection also requires disjoint calibration examples. The post-softmax swap is an off-manifold counterfactual. Preserving the full weight spectrum isolates assignment, but does not show that training would produce the same row or that the intervention is a practical inference-time policy.

Our empirical scope covers controlled retrieval and four pretrained families, including one quantized 7B replication. Natural QA is causal at fixed context, while divergent length-ladder trajectories show that the circuit does not explain every long-context failure. Higher arity remains competence-limited, the generation audit is short-horizon, and utility ranks swaps better than it calibrates large effects.

References

- Brunner, G.; Liu, Y.; Pascual, D.; Richter, O.; Ciaramita, M.; and Wattenhofer, R. 2020. On Identifiability in Transformers. In *International Conference on Learning Representations*.
- Conmy, A.; Mavor-Parker, A.; Lynch, A.; Heimersheim, S.; and Garriga-Alonso, A. 2023. Towards Automated Circuit Discovery for Mechanistic Interpretability. In *Advances in Neural Information Processing Systems*, volume 36.
- Cunningham, H. J.; and Muca Cirone, N. 2026. Contribution Weights: A Geometrical Analysis of Self-Attention Transformers. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*.
- Geva, M.; Bastings, J.; Filippova, K.; and Globerson, A. 2023. Dissecting Recall of Factual Associations in Auto-Regressive Language Models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 12216–12235. Association for Computational Linguistics.
- Hassid, M.; Peng, H.; Rotem, D.; Kasai, J.; Montero, I.; Smith, N. A.; and Schwartz, R. 2022. How Much Does Attention Actually Attend? Questioning the Importance of Attention in Pretrained Transformers. In *Findings of the Association for Computational Linguistics: EMNLP 2022*.
- Hu, X.; Teng, Z.; Zhao, J.; Wu, W.; and Tu, K. 2025. Efficient Length-Generalizable Attention via Causal Retrieval for Long-Context Language Modeling. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, 24727–24743. PMLR.
- Jain, S.; and Wallace, B. C. 2019. Attention Is Not Explanation. In *Proceedings of NAACL-HLT*, 3543–3556.
- Kazemnejad, A.; Padhi, I.; Ramamurthy, K. N.; Das, P.; and Reddy, S. 2023. The Impact of Positional Encoding on Length Generalization in Transformers. arXiv:2305.19466.
- Kobayashi, G.; Kuribayashi, T.; Yokoi, S.; and Inui, K. 2020. Attention Is Not Only a Weight: Analyzing Transformers with Vector Norms. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, 7057–7075. Association for Computational Linguistics.
- Kramár, J.; Lieberum, T.; Shah, R.; and Nanda, N. 2024. AtP*: An Efficient and Scalable Method for Localizing LLM Behaviour to Components. arXiv:2403.00745.
- Li, R.; Boduljak, G.; and Zhou, J. 2025. On Vanishing Variance in Transformer Length Generalization. arXiv:2504.02827.
- Naim, O.; and Asher, N. 2024. On Explaining with Attention Matrices. In *ECAI 2024*, 1035–1042. IOS Press.
- Nakanishi, K. M. 2025. Scalable-Softmax Is Superior for Attention. arXiv:2501.19399.
- Syed, A.; Rager, C.; and Conmy, A. 2023. Attribution Patching Outperforms Automated Circuit Discovery. arXiv:2310.10348.
- Tarzanagh, D. A.; Li, Y.; Zhang, X.; and Oymak, S. 2023. Max-Margin Token Selection in Attention Mechanism. In *Advances in Neural Information Processing Systems*, volume 36.
- Vasylenko, P.; Pitorro, H.; Martins, A. F. T.; and Treviso, M. 2025. Long-Context Generalization with Sparse Attention. arXiv:2506.16640.
- Wang, J.; Ji, T.; Wu, Y.; Yan, H.; Gui, T.; Zhang, Q.; Huang, X.; and Wang, X. 2024. Length Generalization of Causal Transformers without Position Encoding. arXiv:2404.12224.
- Wiegrefe, S.; and Pinter, Y. 2019. Attention Is Not Not Explanation. In *Proceedings of EMNLP-IJCNLP*, 11–20.
- Yamakoshi, T.; McClelland, J. L.; Goldberg, A. E.; and Hawkins, R. D. 2023. Causal Interventions Expose Implicit Situation Models for Commonsense Language Understanding. In *Findings of the Association for Computational Linguistics: ACL 2023*, 13265–13293.
- Zsámboki, P.; Levi, B.; Smith, D. A. J.; Kagalwala, M.; Kell, A.; Liechty, S.; and Wang, C. 2025. Post-Norm can Resharpen Attention. arXiv:2510.08341.