
Attention Concentration Does Not Identify Routing: Spectrum-Preserving Swaps Separate Capacity from Assignment

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Attention concentration cannot identify where a transformer routes information:
2 two attention rows may share the same weight spectrum while assigning their
3 largest weights to different tokens. This is a measurement validity problem. En-
4 tropy, the maximum weight, every norm, and the sorted spectrum are invariant to
5 token permutation, so none of them can vary with the quantity a routing claim is
6 about, and none can serve as evidence for it. We introduce the spectrum-preserving
7 swap (SPS), a causal intervention that preserves every permutation-invariant prop-
8 erty of an attention row while changing its assignment to task-labeled positions.
9 Across controlled retrieval, pretrained language models, and context-grounded
10 question answering, the intervention separates selective capacity from correct as-
11 signment and reveals assignment-responsive, saturated, and selector-mismatch
12 regimes. A local value-aware utility term ranks finite margin effects in controlled
13 retrieval with 0.989 Spearman correlation and 99.5% sign agreement, and utility-
14 selected circuits yield positive held-out effects on natural question answering.
15 Context-length sweeps show where these diagnoses stop being sufficient: in some
16 cases, behavior degrades without corresponding routing changes. SPS gives a con-
17 trolled test of whether attention assignment causally affects a model’s prediction.

18 1 Introduction

19 Interpretability claims about transformers often rest on where attention concentrates. A sharper row,
20 a lower entropy, or a larger maximum weight is read as evidence that a model is attending to the right
21 evidence. We ask whether that inference is licensed. At fixed concentration, does assigning the same
22 attention spectrum to task-required evidence change the prediction?

23 The inference is not licensed, and the reason is structural. Entropy, the maximum weight, every norm,
24 and the full sorted spectrum are invariant to permutations of token positions. Two attention rows can
25 therefore share all of these statistics while one places its largest weight on the required evidence and
26 the other places it on a distractor. A measurement that takes the same value on both rows cannot be
27 evidence for either, so concentration is not an identifying measure of task-correct routing.

28 The question matters where routing accounts are used. When performance declines with length, a
29 transformer may learn a retrieval rule on short sequences and then lose it as distractors accumulate.
30 We see this pattern in copying, key-value recall, and algorithmic tasks [Kazemnejad et al., 2023,
31 Wang et al., 2024, Vasylenko et al., 2025]. As contexts grow, attention often becomes more diffuse,
32 entropy rises, and attention-output variance can collapse [Li et al., 2025]. These trends describe scale
33 and concentration, yet they do not tell us where the attention goes.




Figure 1: Question-answering use case. For the same context and question, a spectrum-preserving swap (SPS) moves the maximum attention weight from a distractor (“Paris”) to required evidence (“Tokyo”). The model, values, and weight spectrum remain unchanged; the paired target-answer margin tests whether assignment affects the answer.

We test this with a spectrum-preserving swap (SPS). At the answer query, we exchange the weight on a known source with a distractor weight and then run the unchanged model. Because the exchange preserves the complete spectrum, a distractor-only exchange gives a matched control. This paired design also lets us ask when moving weight to the source helps. The answer depends on the spectrum’s selective capacity, its assignment to task-labeled positions, and the downstream utility of the moved values.

Source-max raises per-token accuracy by 0.153, source-min lowers it by 0.043, and the distractor control changes it by 0.011. As the spectrum sharpens, the gap between the two assignments grows even though sharpening the original assignment remains near zero, and a local utility term orders the finite margin effect. Qwen and Pythia respond in the same direction, while Gemma and SmoLLM2 expose saturation and selector-mismatch regimes.

Our contributions are:

- **Fixed-spectrum identification.** We separate selective capacity from task-conditioned assignment by preserving the complete attention spectrum and altering only its token placement.
- **Utility-conditioned mechanism.** We characterize when source-donor reassignment changes the final output and test the prediction against exact interventions.
- **Transfer and boundaries.** We test the intervention across controlled tasks, pretrained models, multi-evidence computation, and natural QA, and identify saturation and selector-mismatch regimes.

2 Related Work

Attention concentration and what it does not identify. Attention concentration has been connected to optimization, token selection, extrapolation, and sparse long-context computation [Tarzanagh et al., 2023, Wang et al., 2024, Nakanishi, 2025, Vasylenko et al., 2025]. Other accounts emphasize attention-output scale or normalization [Li et al., 2025, Zsámboki et al., 2025]. These statistics often move together as context length changes. They are used as evidence in routing accounts, but their status as identifying measures of task-correct routing is generally untested. Our intervention tests that identification question while the entire spectrum is held fixed.

Attention intervention and non-identifiability. Attention weights have long been studied as explanations, perturbed to test prediction sensitivity, and shown to be non-identifiable without their value vectors [Jain and Wallace, 2019, Wiegrefe and Pinter, 2019, Brunner et al., 2020, Kobayashi et al., 2020, Hassid et al., 2022]. Causal circuit analysis patches internal activations and measures the resulting output change [Yamakoshi et al., 2023, Conmy et al., 2023]. Attribution patching uses a gradient approximation to screen many candidate components efficiently [Syed et al., 2023, Kramár et al., 2024]. Our utility term has a related first-order form. Here we evaluate it along an exact permutation of a model’s own attention row, alongside the finite swap itself.

70 **Closest attention interventions.** The nearest methods delete, block, renormalize, or rank parts of
 71 an attention pathway: Attention Knockout blocks selected edges and renormalizes the row [Geva et al.,
 72 2023], efficient attention projects a row onto its output-effective subspace [Naim and Asher, 2024],
 73 and Contribution Weights rank tokens by attention mass combined with projected-value magnitude,
 74 testing importance through removal [Cunningham and Muca Cirone, 2026]. Each changes the weight
 75 multiset. Our reassignment leaves it intact and changes only its assignment to token positions, and
 76 Equation 5 connects that finite, task-labeled swap to the target margin. Sparse and retrieval-based
 77 architectures instead change how distant context is accessed [Vasylenko et al., 2025, Hu et al., 2025];
 78 we study a trained row with a fixed spectrum and isolate the effect of assigning its available selectivity
 79 to task-required evidence.

80 3 Methods

81 3.1 Capacity and Assignment

82 At an answer query, let $a = (a_1, \dots, a_n)$ be an attention row over n token positions, where $a_j \geq 0$
 83 and $\sum_j a_j = 1$. Let $S \subseteq \{1, \dots, n\}$ contain the $r = |S|$ task-required evidence positions, and let
 84 $a_{(1)} \geq \dots \geq a_{(n)}$ denote the same weights in descending order. Define

$$C_r^+ = \sum_{i=1}^r a_{(i)}, \quad C_r^- = \sum_{i=n-r+1}^n a_{(i)}, \quad (1)$$

$$K_r(a) = C_r^+ - C_r^-.$$

85 C_r^+ and C_r^- are the largest and smallest evidence masses available under permutations of the fixed
 86 spectrum, and $K_r(a)$ is the available routing range. An assignment is a permutation π of $\{1, \dots, n\}$
 87 that places weight $a_{\pi(j)}$ at token position j . Its evidence mass and normalized assignment coordinate
 88 are

$$m_S(\pi) = \sum_{j \in S} a_{\pi(j)}, \quad (2)$$

$$\rho_S(\pi) = \frac{m_S(\pi) - C_r^-}{K_r(a)} \in [0, 1], \quad K_r(a) > 0.$$

89 Thus $m_S(\pi) = C_r^- + \rho_S(\pi)K_r(a)$. If $K_r(a) = 0$, every assignment has the same evidence mass.

90 **Capacity** specifies how much evidence mass the spectrum can support.
Assignment specifies where that mass is placed.
Utility specifies whether moving it changes the answer.

91 **Proposition 1** (Fixed-spectrum capacity-assignment decomposition). *For every assignment π of a*
 92 *fixed weight multiset,*

$$C_r^- \leq m_S(\pi) \leq C_r^+, \quad (3)$$

93 *and both bounds are attained. If $K_r(a) > 0$, Equation 2 defines a unique $\rho_S(\pi)$. No statistic*
 94 *invariant to token permutation can determine ρ_S .*

95 **Proof sketch.** An r -position evidence set receives at most the r largest weights and at least the r
 96 smallest. Assigning those weights to S attains each endpoint. Normalizing evidence mass within this
 97 interval gives ρ_S . Any permutation-invariant statistic is constant across all assignments, including
 98 assignments at opposite endpoints. The supplementary material gives the full exchange argument
 99 and the corresponding affine margin range. The decomposition is algebraic for every r . Behavioral
 100 validation is positive for $r = 1$ and $r = 2$; the $r = 3$ interval includes zero and the $r = 4$ models
 101 miss competence, so efficacy at higher arity is unestablished.

102 3.2 Utility-Conditioned Output Effects

103 Capacity and assignment describe attention mass. Prediction also depends on values and downstream
 104 computation. Let H be the selected attention heads. In each $h \in H$, let s be the evidence position,
 105 d_h the donor position whose weight is exchanged with it, and a_{hj} the original attention weight on
 106 position j . Define the transferred mass $\delta_h = a_{hd_h} - a_{hs}$. Let x_{hj} be token j 's residual-stream

107 contribution after head h 's value and output projections. If z is the unpatched residual state, the
 108 intervention gives $z' = z + \Delta z$, where

$$\Delta z = \sum_{h \in H} \delta_h (x_{hs} - x_{hd_h}). \quad (4)$$

109 Let $\ell_k(z)$ be the downstream logit for token k , with target y and competitor c . Define $M_c(z) =$
 110 $\ell_y(z) - \ell_c(z)$, the exact change $\Delta M_c = M_c(z + \Delta z) - M_c(z)$, and local utility $u_{hj}^{(c)} =$
 111 $\nabla_z M_c(z)^\top x_{hj}$.

112 **Proposition 2** (Utility-conditioned swap effect). *If $\nabla_z M_c$ is L_c -Lipschitz on $\{z + t\Delta z : t \in [0, 1]\}$,*
 113 *then*

$$\Delta M_c = \sum_{h \in H} \delta_h (u_{hs}^{(c)} - u_{hd_h}^{(c)}) + R, \quad |R| \leq \frac{L_c}{2} \|\Delta z\|_2^2. \quad (5)$$

114 For an affine downstream margin, $R = 0$.

115 The Taylor remainder is standard. The new component is the task-labeled capacity-assignment
 116 coordinate coupled to the exact SPS vector in Equation 4.

117 The intervention transfers attention mass from a donor token to the source. It helps when the
 118 source contributes more positively to the target margin than the displaced donor, unless downstream
 119 curvature overwhelms that gain. Source-max can therefore succeed, saturate, or fail under the same
 120 attention-mass rule. The supplementary material derives Equation 5, extends it to evidence sets, and
 121 states a sufficient margin certificate.

122 4 Experimental Setup

123 **Controlled tasks.** Argmax retrieval presents key-value pairs and requests the value paired with
 124 the largest key. Flag retrieval presents values with exactly one marked position and requests that
 125 value. Both have one labeled source. Two-evidence modular sum marks two digits and requests
 126 their sum modulo ten, so both sources are required. Four-layer decoders are trained with NoPE and
 127 RoPE and four seeds per cell; evaluation reaches $50\times$ training length. An eight-layer, width-512
 128 grid checks scale. The inverse temperature $\beta \in \{1, 2, 4\}$ changes attention capacity while each
 129 assignment comparison preserves the resulting sorted spectrum.

130 **Pretrained and natural-language tasks.** In-context key-value recall lists pairs followed by a query
 131 key; the paired value is the labeled source. We test Qwen2.5, Pythia, Gemma, and SmoLLM2 across
 132 increasing pair counts. The selector audit compares circuits ranked on disjoint calibration data by
 133 source mass or by output-conditioned utility. Natural MCQA uses four-choice questions with one
 134 gold passage and unrelated passages. Questions must be correct with the full context and gold passage
 135 alone, but incorrect without context. This filtering defines a model-specific context-dependent subset
 136 and is not intended to estimate average performance over the full dataset. Its length ladder freezes the
 137 four-passage circuit before forming nested 8-, 16-, and 32-passage contexts with additional unrelated
 138 passages.

139 **Interventions and estimation.** The distractor control swaps the largest and smallest non-source
 140 weights, preserving source mass while permuting the same attention row. For a source position s , let
 141 $d^+ = \arg \max_j a_j$, $d^- = \arg \min_j a_j$, and $D = \{j : j \neq s\}$ over valid tokens. Source-max swaps
 142 (s, d^+) and source-min swaps (s, d^-) . In controlled and pretrained key-value recall, the distractor
 143 control uses the maximum and minimum within D . Natural QA uses a displacement-matched control.
 144 With $\delta = a_{d^+} - a_s$, it chooses

$$(p, q) = \arg \min_{\{j, k\} \subset D} |a_j - a_k| - \delta \quad (6)$$

145 and swaps (p, q) . Among all non-source pairs, this choice minimizes the absolute difference between
 146 the control's L_1 displacement and the treatment's L_1 displacement; the remaining difference is
 147 the minimum attainable mismatch for that row. It holds source mass fixed, and every control
 148 preserves the complete spectrum. For natural QA, displacement is matched first because it measures
 149 perturbation magnitude. The control does not match the donor token, positional distance, or value

norm, which remain separate confounds. Calibration and evaluation examples are disjoint, circuits are frozen before paired evaluation, and hierarchical intervals resample calibration or model seeds above paired examples. We interpret a cell only after its frozen competence gate passes. Three confirmatory contrasts form one inferential family: controlled source-max versus distractor control, the capacity-by-assignment interaction, and natural utility-selected source-max versus matched control. We use two-sided exact sign-flip tests at the trained-model or calibration-seed cluster level and Holm correction across the family. Other intervals are descriptive and unadjusted. These swaps are counterfactual probes and need not lie on the model’s endogenous query-key manifold. Treatment and control share this intervention class, so their contrast measures downstream sensitivity to assignment, rather than the frequency with which the model naturally creates a particular row. Layer controls (+1.107 for selected $K = 8$ versus -0.059 for a random layer) and coverage dose response show greater sensitivity under this probe. They do not show that normal query-key computation produces swapped rows or that selected heads are the exclusive retrieval mechanism. In untouched distractor-order variants, the conditional source-mass coefficient is +1.250 ($p = .03125$), while the closest-spectrum paired effect is +.026 ($[-.017, +.070]$, $p = .25$). This connects the intervention to an endogenous association without establishing endogenous causality.




Figure 2: Paired causal protocol. A circuit selected on disjoint calibration data is frozen before treatment and control are applied to the same evaluation row. Both paths preserve the spectrum; the natural QA control additionally minimizes displacement mismatch in Equation 6.

Measured outcomes. We record task accuracy, target-versus-competitor margin, source mass, attention-output variance, entropy, norms, and the full sorted spectrum. Every fixed-spectrum condition preserves the selected attention rows to numerical precision. The supplementary material reports prompts, hyperparameters, per-length tables, competence failures, and preregistered decision rules.




Figure 3: Assignment changes retrieval despite an identical attention spectrum. Source-max improves accuracy, source-min harms it, and the distractor-only control remains near zero (a). The effect increases with the fraction of the selected circuit receiving the intervention (b). Diamonds and intervals summarize trained models; faint paths show individual cells.

5 Results and Discussion

5.1 Assignment Changes Accuracy at Fixed Spectrum

We begin with the simplest test: move the largest available weight to the source, move it away, or exchange two distractors. Across sixteen trained models, the mean changes are +0.153, -0.043 , and $+0.011$, respectively (Figure 3). Source-max beats the control in every model-length comparison. Their primary cluster-level contrast is $+0.142$ (95% CI $[+0.080, +0.206]$; exact $p = 3.05 \times 10^{-5}$; Holm-adjusted $p = 9.16 \times 10^{-5}$), while complete-spectrum invariant error stays below 10^{-6} .

The controlled audit contains 8,192 head-example pairs. Individual swaps are vacuous in 18.5% of them, but only 86 of 1,024 complete eight-head circuits (8.4%) are entirely vacuous. Mean exact margin change is $+3.891$ over all circuits and $+4.248$ (95% CI $[+3.863, +4.640]$) over the 938 active circuits. The median per-head transferred mass is .025 (90th percentile .930), and effect size rises across transfer quartiles. The supplementary displacement audit reports the full distribution.

The effect grows with circuit coverage (Figure 3b). Patching one, two, four, and eight selected heads gives mean source-max gains of 0.007, 0.013, 0.056, and 0.153 and is monotone in 31 of 32 model-length comparisons, consistent with a distributed effect under the probe.

Source attention tracks accuracy more closely than attention-output variance. Normalizing that variance leaves accuracy unchanged, whereas forcing mass onto the known source restores controlled retrieval. Sharpening the original row has little effect.

The ordering survives the eight-layer replication, both positional encodings, four training seeds, and two controlled tasks. Random-layer and adjacent-layer interventions remain near zero when the selected circuit is held fixed. Full model-level intervals and the variance and patching results appear in the supplementary material.

5.2 Capacity Amplifies Assignment

We next vary spectrum and assignment independently. Temperature scaling changes the maximum available weight, and at every temperature we evaluate identity, source-max, source-min, and distractor control on the same examples. The source-max minus source-min contrast grows from 0.196 to 0.346, while sharpening the identity assignment changes accuracy by at most 0.009 (Figure 4a). Capacity sets the possible leverage, while assignment determines where that leverage lands. The primary interaction is $+0.150$ (95% CI $[+0.094, +0.210]$; exact $p = 6.10 \times 10^{-5}$; Holm-adjusted $p = 1.22 \times 10^{-4}$).

For held-out examples we differentiate the target margin along the exact source-max path. The first-order term in Proposition 2 predicts the finite margin change with 0.917 Pearson correlation, 0.989 Spearman correlation, and 99.5% sign agreement (Figure 4b). The remaining error measures downstream curvature along the swap.




Figure 4: Capacity controls leverage, and utility predicts the finite intervention. Sharper spectra increase the correct-versus-wrong assignment contrast while sharpening alone remains near zero (a). The first-order utility term predicts the exact source-max margin effect across tasks and positional encodings (b). Error bars summarize trained models.




Figure 5: Matched pretrained-model comparison (seed 0; $n = 128$ paired examples per model-length cell). All panels use the 5-, 20-, and 80-pair ladder and report source-max, source-min, and untouched margin relative to the same distractor control. Qwen and Pythia respond positively, Gemma is saturated, and the source-mass-selected SmoLLM2 circuit does not benefit. Error bars show paired 95% intervals.

205 Modular sum requires the model to route and combine two marked values. At $5\times$ training length,
 206 assigning the top two weights to evidence raises exact match by 0.066, assigning the bottom two
 207 lowers it by 0.291, and evidence-max exceeds a control that preserves both the spectrum and total
 208 evidence mass by 0.043 (95% CI [0.017, 0.072]). The result therefore extends beyond copying one
 209 token. A registered three-evidence estimate includes zero, and four-evidence models miss the frozen
 210 competence gate. Higher arity remains unresolved.

211 Evidence-max and evidence-min preserve the complete spectrum. The multi-evidence control also
 212 preserves total evidence mass, changing only which evidence token receives each weight. Its small
 213 effect separates aggregate source mass from useful evidence assignment.

214 5.3 The Same Intervention Produces Three Pretrained Regimes

215 The pretrained models show why the utility term matters. Qwen and Pythia gain margin when
 216 source-max replaces the distractor control. Gemma stays accurate as context grows and gains little
 217 from extra source mass, even though source-min reliably harms it. This is the saturation regime
 218 predicted by Proposition 2: removing evidence matters, while extra routing has little left to add. A
 219 quantized Qwen2.5-7B run preserves the positive margin direction.

220 SmoLLM2 provides the contrasting case. Its preregistered source-mass selector fails, yet ranking
 221 heads on disjoint calibration examples by output-conditioned gradient or utility restores a positive

Fixed-spectrum assignment		n Units	Unpatched	Δ source-max	Δ source-min	Δ distractor	max – distractor
Single evidence	$\beta = 1$	8,192 Acc.	.619	+.153	-.043	+.011	+.142
	$\beta = 4$	8,192 Acc.	.619	+.298	-.048	+.025	+.273
Two evidence	1x	2,048 EM	.952	+.019	-.559	+.012	+.007
	5x	2,048 EM	.484	+.066	-.291	+.023	+.043
Pretrained KV	Qwen2.5-1.5B	384 Margin	-3.048	+1.990	-.961	+.192	+1.799
	Pythia-1.4B	384 Margin	-3.805	+.826	-1.022	+.146	+.680
	Gemma-2-2B	384 Margin	+2.234	+.040	-1.373	+.013	+.027
	SmolLM2-1.7B	384 Margin	-2.782	-.168	-.112	-.077	-.091
Output-conditioned circuit selection		n Units	ΔM versus matched control		ΔM source mass	$\Pr(\Delta M > 0)$	utility – source
Task	Model		Utility gain	Utility gap			
Pretrained KV	Pythia-1.4B	384 Margin	+.519	+.359	+.006	.810	+.513
	SmolLM2-1.7B	384 Margin	+1.253	+.859	+.006	.969	+1.247
	Qwen2.5-1.5B	384 Margin	+1.493	+1.493	+.069	.956	+1.424

Table 1: Benchmark summary. The pretrained fixed-spectrum rows use the matched 5-, 20-, and 80-pair ladder from Figure 5, with 128 paired examples per model-length cell and $n = 384$ per model. Units make upper-block baselines comparable only within task. The lower block reports the three model families with completed matched selector runs; a Gemma selector-family run was not conducted. It has no unpatched baseline and evaluates selector quality using fixed-spectrum control effects, not per-row utility prediction; the supplementary held-out audit evaluates that prediction. $\Pr(\Delta M > 0)$ is the fraction of paired evaluation examples with a positive utility-gain source-max minus matched-control margin. Bold final columns remain identifiable in grayscale. Full cells and 95% intervals are in the supplementary material.




Figure 6: Matched natural-QA length ladder (seed 0; $n = 128$ paired examples per model-length cell). From 4 to 32 passages, accuracy falls (a), source routing moves in opposite directions (b), and frozen-circuit rescue stays flat or weakens (c). Bands are paired-example bootstrap 95% intervals.

effect. A ten-seed ablation ranks source gradient, utility gain, and utility gap above source mass and transfer mass. Rankings vary by architecture, but circuit selection must account for the output.

Interpolating from the natural row to the utility-selected swap gives a monotone margin path over five calibration seeds. On context-grounded natural QA, utility-selected routing improves margin over a displacement-matched control, while source-mass selection harms it. Across six disjoint calibration seeds, the utility-selected contrast is $+.378$ (hierarchical 95% CI $[+.253, +.520]$; exact and Holm-adjusted $p = .03125$). All competence gates pass. An unconstrained generation audit finds no first-answer or repetition cost.

The selector results judge a selector by the frozen circuits it finds and their paired margin effects. They do not test whether the per-row utility functional predicts a finite swap effect. That held-out predictor audit is reported separately in the supplementary material.

Clean-to-corrupt activation patching rescues $+.308$ margin but only between 2.5% and 3.3% of median corruption damage. It restores selected-head outputs after source-value corruption, so its magnitude is not directly comparable to reassignment on an unchanged input.

5.4 Natural Context Marks the Scope Boundary

The fixed-context natural-QA result establishes causal routing in ordinary text. The length ladder is a scope test: it asks whether the same fixed-context protocol remains informative as context grows. We select the circuit at four passages, freeze every choice, and append unrelated passages to form nested 4, 8, 16, and 32-passage contexts. The matched figure fixes seed 0 and resamples paired questions; repeated passage observations from one circuit are not treated as independent models. The supplementary material reports the additional Qwen seed separately.

Both Qwen and SmoLLM2 lose margin and accuracy by 32 passages, yet their source-mass trajectories diverge. Qwen source mass rises as intervention rescue weakens. SmoLLM2 source mass falls while rescue stays flat. The shared degradation cannot be attributed to one consistent source-mass trajectory. This is a scope limit of the diagnosis, not a failed prediction.

The protocol has a narrower interpretation. At a fixed context, a spectrum-preserving swap measures whether a chosen circuit can use reassigned evidence. As contexts grow, representation quality, distractor semantics, answer competition, and computation outside that circuit can also matter. The theorem concerns a specified swap at a fixed context; this experiment does not resolve the other long-context bottlenecks. Exact 32-minus-4 changes and hierarchical intervals appear in the supplementary material.

The arity checks provide another limit. The two-evidence task passes and supports the mass-preserving control; the three-evidence interval includes zero, and the four-evidence model misses competence. Higher arity remains unresolved. An eight-token generation audit also leaves first-answer accuracy and repetition unchanged, so simple format collapse is not driving the result.

5.5 Identification and Replication Audits

Replication checks vary the calibration sample, patch budget, and model family, and they support the same reading. Selected circuits are positive at every head budget, adjacent and random layers are smaller or near zero, the utility term ranks exact effects within every competent cell while calibrating them imperfectly, and an exact rerun reproduces the direction with an atypically large original magnitude. Appendix A reports these audits in full.

6 Limitations

SPS identifies the causal effect of a specified attention reassignment, not a complete account of model computation. The protocol requires identifiable evidence positions and a chosen intervention circuit; output-conditioned selection also requires disjoint calibration examples. The post-softmax swap is an off-manifold counterfactual. Preserving the full weight spectrum isolates assignment, but does not show that training would produce the same row or that the intervention is a practical inference-time policy.

Our empirical scope covers controlled retrieval and four pretrained families, including one quantized 7B replication. Natural QA is causal at fixed context, while divergent length-ladder trajectories show that the circuit does not explain every long-context failure. Higher arity remains competence-limited, the generation audit is short-horizon, and utility ranks swaps better than it calibrates large effects.

References

- Gino Brunner, Yang Liu, Damián Pascual, Oliver Richter, Massimiliano Ciaramita, and Roger Wattenhofer. On identifiability in transformers. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=BJg1f6EFDB>.
- Arthur Conmy, Augustine Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/34e1dbe95d34d7ebaf99b9bcaeb5b2be-Abstract-Conference.html.
- Harry Jake Cunningham and Nicola Muca Cirone. Contribution weights: A geometrical analysis of self-attention transformers. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026. URL <https://arxiv.org/abs/2606.07604>.

285 Mor Geva, Jasmijn Bastings, Katja Filippova, and Amir Globerson. Dissecting recall of factual associations
286 in auto-regressive language models. In *Proceedings of the 2023 Conference on Empirical Methods in*
287 *Natural Language Processing*, pages 12216–12235. Association for Computational Linguistics, 2023. doi:
288 10.18653/v1/2023.emnlp-main.751. URL <https://aclanthology.org/2023.emnlp-main.751/>.

289 Michael Hassid, Hao Peng, Daniel Rotem, Jungo Kasai, Ivan Montero, Noah A. Smith, and Roy Schwartz.
290 How much does attention actually attend? questioning the importance of attention in pretrained trans-
291 formers. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, 2022. URL
292 <https://aclanthology.org/2022.findings-emnlp.101/>.

293 Xiang Hu, Zhihao Teng, Jun Zhao, Wei Wu, and Kewei Tu. Efficient length-generalizable attention via causal
294 retrieval for long-context language modeling. In *Proceedings of the 42nd International Conference on*
295 *Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 24727–24743. PMLR,
296 2025. URL <https://proceedings.mlr.press/v267/hu25v.html>.

297 Sarthak Jain and Byron C. Wallace. Attention is not explanation. In *Proceedings of NAACL-HLT*, pages
298 3543–3556, 2019. URL <https://aclanthology.org/N19-1357/>.

299 Amirhossein Kazemnejad, Inkit Padhi, Karthikeyan Natesan Ramamurthy, Payel Das, and Siva Reddy. The
300 impact of positional encoding on length generalization in transformers, 2023.

301 Goro Kobayashi, Tatsuki Kuribayashi, Sho Yokoi, and Kentaro Inui. Attention Is Not Only a Weight: Analyzing
302 transformers with vector norms. In *Proceedings of the 2020 Conference on Empirical Methods in Natural*
303 *Language Processing*, pages 7057–7075. Association for Computational Linguistics, 2020. doi: 10.18653/v1/
304 2020.emnlp-main.574. URL <https://aclanthology.org/2020.emnlp-main.574/>.

305 János Kramár, Tom Lieberum, Rohin Shah, and Neel Nanda. AtP*: An efficient and scalable method for
306 localizing LLM behaviour to components, 2024. URL <https://arxiv.org/abs/2403.00745>.

307 Ruining Li, Gabor Boduljak, and Jinghao Zhou. On vanishing variance in transformer length generalization,
308 2025.

309 Omar Naim and Nicholas Asher. On explaining with attention matrices. In *ECAI 2024*, pages 1035–1042. IOS
310 Press, 2024. doi: 10.3233/FAIA240594. URL <https://doi.org/10.3233/FAIA240594>.

311 Ken M. Nakanishi. Scalable-softmax is superior for attention, 2025.

312 Aaqib Syed, Can Rager, and Arthur Conmy. Attribution patching outperforms automated circuit discovery,
313 2023. URL <https://arxiv.org/abs/2310.10348>.

314 Davoud Ataee Tarzanagh, Yingcong Li, Xuechen Zhang, and Samet Oymak. Max-margin to-
315 ken selection in attention mechanism. In *Advances in Neural Information Processing Systems*,
316 volume 36, 2023. URL [https://proceedings.neurips.cc/paper_files/paper/2023/hash/](https://proceedings.neurips.cc/paper_files/paper/2023/hash/970f59b22f4c72aec75174aae63c7459-Abstract-Conference.html)
317 [970f59b22f4c72aec75174aae63c7459-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/970f59b22f4c72aec75174aae63c7459-Abstract-Conference.html).

318 Pavlo Vasylenko, Hugo Pitorro, André F. T. Martins, and Marcos Treviso. Long-context generalization with
319 sparse attention, 2025.

320 Jie Wang, Tao Ji, Yuanbin Wu, Hang Yan, Tao Gui, Qi Zhang, Xuanjing Huang, and Xiaoling Wang. Length
321 generalization of causal transformers without position encoding, 2024.

322 Sarah Wiegrefe and Yuval Pinter. Attention is not not explanation. In *Proceedings of EMNLP-IJCNLP*, pages
323 11–20, 2019. URL <https://aclanthology.org/D19-1002/>.

324 Takateru Yamakoshi, James L. McClelland, Adele E. Goldberg, and Robert D. Hawkins. Causal interventions
325 expose implicit situation models for commonsense language understanding. In *Findings of the Association for*
326 *Computational Linguistics: ACL 2023*, pages 13265–13293, 2023. doi: 10.18653/v1/2023.findings-acl.839.
327 URL <https://aclanthology.org/2023.findings-acl.839/>.

328 Pál Zsámboki, Benjamin Levi, David Ansel Josef Smith, Mitansh Kagalwala, Arlington Kell, Samuel Liechty,
329 and Cong Wang. Post-norm can sharpen attention, 2025.

330 **A Identification and replication audits**

331 Replication checks vary the calibration sample, patch budget, and model family. Across three frozen
332 Qwen calibration seeds, selected circuits are positive at every head budget: $K = 2, 4, 8$ gives $+ .460$,
333 $+1.054$, and $+1.107$ margin. Adjacent layers are smaller and inconsistent ($+ .147$ and $- .131$), while
334 random layers average near zero ($- .059$). Full intervals and seed standard deviations are reported in
335 the supplementary package.

336 A separate replication tests the utility relation. Circuits are chosen by source mass on disjoint
337 calibration examples; utility is then computed on held-out evaluation rows. Across three seeds per
338 family, every competent source-max cell is positive, and predicted effects rank the exact effects
339 within each cell. Mean Spearman correlations are $.866$ for Qwen2.5-1.5B, $.544$ for Pythia-1.4B, and
340 $.554$ for Gemma-2-2B, with sign agreements $.805$, $.807$, and $.714$. The approximation is imperfect at
341 large effects, so we use it as a local ranking signal rather than a calibrated finite-effect estimator. For
342 the 32 controlled cell means in Figure 4b, the actual-on-predicted fit is $1.387x - .102$ (Pearson $.978$),
343 which deviates from identity calibration.

344 The exact Qwen rerun keeps precision, lengths, contrast, and sample counts fixed. All five new seeds
345 are positive (mean $+ .042$); the combined six-seed test gives $p = .03125$, while the original $+1.882$
346 magnitude is atypical.

347 Output-conditioned selectors lead in each family, although their internal ranking changes by architec-
348 ture. Full selector results, competence-gate failures, and the anonymized code package appear in the
349 supplementary material.