
Fixed-Spectrum Swaps Separate Attention Capacity from Assignment

Anonymous Author(s)

Affiliation

Address

email

Abstract

Interpretability often reads attention mass on a gold token as evidence that a transformer routes information to it. We show that this inference is unsafe in a specific, testable way. Statistics of the attention weights alone, including entropy, the maximum, every norm, and the sorted spectrum, are invariant to which token holds which weight, so they cannot identify assignment; and attention mass on the gold token, which does identify assignment, still does not determine the causal effect of that assignment on the output. We introduce the fixed-spectrum swap (FSS), a causal intervention that preserves every permutation-invariant property of an attention row while changing its assignment to task-labeled positions. Across controlled retrieval, pretrained language models, and context-grounded question answering, FSS separates selective capacity from correct assignment and reveals assignment-responsive, saturated, and selector-mismatch regimes. Outcome-blind selection by source mass fails where outcome-aware selection succeeds; because our utility term is the gradient of the margin it predicts, this establishes that outcome-aware selection helps rather than isolating a special role for utility, and separating the two needs an outcome-aware but non-utility baseline. Pretrained effect magnitudes are seed-sensitive, so we report them as directional: on the one condition we replicate across seeds, the original seed is $45\times$ the five-seed mean. On natural question answering the prediction is sensitive to reassignment onto the gold evidence, though a semantic-distractor control is still needed to separate routing to required evidence from routing to any answer-shaped token. FSS gives a controlled test of whether a specified post-softmax reassignment of attention changes a model’s prediction.

1 Introduction

Interpretability claims about transformers often rest on how attention is distributed. Attention weights are read as an explanation of which tokens a prediction depends on [Jain and Wallace, 2019, Wiegrefe and Pinter, 2019], and attention mass on a gold token, or a summary of it such as low entropy or a large maximum, is taken as evidence that the model routes information to that token. We ask whether that inference is licensed, and we separate two versions of it. Summary statistics of the weights alone cannot even identify assignment; attention mass on the gold token can, yet still may not tell us whether the routing causally drives the prediction.

The first gap is structural. Entropy, the maximum weight, every norm, and the full sorted spectrum are invariant to permutations of token positions, so two attention rows can share all of them while one places its largest weight on the required evidence and the other on a distractor. A measurement equal on both rows cannot distinguish them. Attention mass on the gold token, m_S , is the sharper and more common measure: it is not permutation-invariant and does pin down the assignment. Our question is




Figure 1: Question-answering use case. For the same context and question, a fixed-spectrum swap (FSS) moves the maximum attention weight from a distractor (“Paris”) to required evidence (“Tokyo”). The model, values, and weight spectrum remain unchanged; the paired target-answer margin tests whether assignment affects the answer.

therefore not about m_S as a descriptor but about its causal warrant. At a fixed spectrum, does moving mass onto the task-required evidence change the prediction, and is m_S enough to say when it will?

We test this with a fixed-spectrum swap (FSS). At the answer query, we exchange the weight on a known source with a distractor weight and then run the unchanged model. Because the exchange preserves the complete spectrum, a distractor-only exchange gives a matched control. The answer depends on the spectrum’s selective capacity, its assignment to task-labeled positions, and the downstream utility of the moved values, and we find that source mass alone does not determine it: selecting circuits by source mass fails where selecting by a value-aware utility term succeeds.

Length is a boundary we test, not a phenomenon we promise to explain. Performance often declines as distractors accumulate [Kazemnejad et al., 2023, Wang et al., 2024, Vasylenko et al., 2025], and attention grows more diffuse while its output variance can collapse [Li et al., 2025]. We use context-length sweeps to ask how far the fixed-spectrum diagnosis remains informative, and we report where it stops rather than claiming it accounts for long-context failure.

Our contributions are:

- **Fixed-spectrum identification.** We separate selective capacity from task-conditioned assignment by preserving the complete attention spectrum and altering only its token placement.
- **Source mass is not sufficient.** We show that source mass identifies assignment but does not determine the causal effect: outcome-blind source-mass selection fails where outcome-aware selection succeeds.
- **Transfer and boundaries.** We test the intervention across controlled tasks, pretrained models, multi-evidence computation, and natural QA, and identify saturation and selector-mismatch regimes together with the seed, endogenous, selector-circularity, and length boundaries where the reading stops being sufficient.

2 Related Work

Attention weights as routing evidence. Attention weights are widely read as evidence of which tokens a prediction uses, a practice examined in the attention-as-explanation debate [Jain and Wallace, 2019, Wiegrefe and Pinter, 2019, Brunner et al., 2020]. Concentration summaries such as entropy, maximum weight, and output variance are connected to optimization, token selection, and extrapolation [Tarzanagh et al., 2023, Wang et al., 2024, Nakanishi, 2025, Vasylenko et al., 2025, Li et al., 2025, Zsámboi et al., 2025], and move together as context length changes. Our target is not these summaries, which few defend as routing measures, but attention mass on the gold token, which is the measure such claims actually rely on. Mass identifies assignment, so the open question is causal: does that assignment drive the output, and is mass enough to predict when. We test this while the entire spectrum is held fixed.

Attention intervention and non-identifiability. Attention weights have long been studied as explanations, perturbed to test prediction sensitivity, and shown to be non-identifiable without their

value vectors [Jain and Wallace, 2019, Wiegrefe and Pinter, 2019, Brunner et al., 2020, Kobayashi et al., 2020, Hassid et al., 2022]. Causal circuit analysis patches internal activations and measures the resulting output change [Yamakoshi et al., 2023, Conmy et al., 2023]. Attribution patching uses a gradient approximation to screen many candidate components efficiently [Syed et al., 2023, Kramár et al., 2024]. Our utility term has a related first-order form. Here we evaluate it along an exact permutation of a model’s own attention row, alongside the finite swap itself.

Closest attention interventions. The nearest methods delete, block, renormalize, or rank parts of an attention pathway: Attention Knockout blocks selected edges and renormalizes the row [Geva et al., 2023], efficient attention projects a row onto its output-effective subspace [Naim and Asher, 2024], and Contribution Weights rank tokens by attention mass combined with projected-value magnitude, testing importance through removal [Cunningham and Muca Cirone, 2026]. Each changes the weight multiset. Our reassignment leaves it intact and changes only its assignment to token positions, and Equation 5 connects that finite, task-labeled swap to the target margin. Sparse and retrieval-based architectures instead change how distant context is accessed [Vasylenko et al., 2025, Hu et al., 2025]; we study a trained row with a fixed spectrum and isolate the effect of assigning its available selectivity to task-required evidence.

3 Methods

3.1 Capacity and Assignment

At an answer query, let $a = (a_1, \dots, a_n)$ be an attention row over n token positions, where $a_j \geq 0$ and $\sum_j a_j = 1$. Let $S \subseteq \{1, \dots, n\}$ contain the $r = |S|$ task-required evidence positions, and let $a_{(1)} \geq \dots \geq a_{(n)}$ denote the same weights in descending order. Define

$$C_r^+ = \sum_{i=1}^r a_{(i)}, \quad C_r^- = \sum_{i=n-r+1}^n a_{(i)}, \quad (1)$$

$$K_r(a) = C_r^+ - C_r^-.$$

C_r^+ and C_r^- are the largest and smallest evidence masses available under permutations of the fixed spectrum, and $K_r(a)$ is the available routing range. An assignment is a permutation π of $\{1, \dots, n\}$ that places weight $a_{\pi(j)}$ at token position j . Its evidence mass and normalized assignment coordinate are

$$m_S(\pi) = \sum_{j \in S} a_{\pi(j)}, \quad (2)$$

$$\rho_S(\pi) = \frac{m_S(\pi) - C_r^-}{K_r(a)} \in [0, 1], \quad K_r(a) > 0.$$

Thus $m_S(\pi) = C_r^- + \rho_S(\pi)K_r(a)$. If $K_r(a) = 0$, every assignment has the same evidence mass.

Capacity specifies how much evidence mass the spectrum can support.
Assignment specifies where that mass is placed.
Utility specifies whether moving it changes the answer.

Proposition 1 (Fixed-spectrum capacity-assignment decomposition). *For every assignment π of a fixed weight multiset,*

$$C_r^- \leq m_S(\pi) \leq C_r^+, \quad (3)$$

and both bounds are attained. If $K_r(a) > 0$, Equation 2 defines a unique $\rho_S(\pi)$. No statistic invariant to token permutation can determine ρ_S .

Proof sketch. An r -position evidence set receives at most the r largest weights and at least the r smallest. Assigning those weights to S attains each endpoint. Normalizing evidence mass within this interval gives ρ_S . Any permutation-invariant statistic is constant across all assignments, including assignments at opposite endpoints. The supplementary material gives the full exchange argument and the corresponding affine margin range. The decomposition is algebraic for every r . Behavioral validation is positive for $r = 1$ and $r = 2$; the $r = 3$ interval includes zero and the $r = 4$ models miss competence, so efficacy at higher arity is unestablished.

3.2 Utility-Conditioned Output Effects

Capacity and assignment describe attention mass. Prediction also depends on values and downstream computation. Let H be the selected attention heads. In each $h \in H$, let s be the evidence position, d_h the donor position whose weight is exchanged with it, and a_{hj} the original attention weight on position j . Define the transferred mass $\delta_h = a_{hd_h} - a_{hs}$. Let x_{hj} be token j 's residual-stream contribution after head h 's value and output projections. If z is the unpatched residual state, the intervention gives $z' = z + \Delta z$, where

$$\Delta z = \sum_{h \in H} \delta_h (x_{hs} - x_{hd_h}). \quad (4)$$

Let $\ell_k(z)$ be the downstream logit for token k , with target y and competitor c . Define $M_c(z) = \ell_y(z) - \ell_c(z)$, the exact change $\Delta M_c = M_c(z + \Delta z) - M_c(z)$, and local utility $u_{hj}^{(c)} = \nabla_z M_c(z)^\top x_{hj}$.

Proposition 2 (Utility-conditioned swap effect). *If $\nabla_z M_c$ is L_c -Lipschitz on $\{z + t\Delta z : t \in [0, 1]\}$, then*

$$\Delta M_c = \sum_{h \in H} \delta_h (u_{hs}^{(c)} - u_{hd_h}^{(c)}) + R, \quad |R| \leq \frac{L_c}{2} \|\Delta z\|_2^2. \quad (5)$$

For an affine downstream margin, $R = 0$.

The Taylor remainder is standard. The new component is the task-labeled capacity-assignment coordinate coupled to the exact FSS vector in Equation 4.

The intervention transfers attention mass from a donor token to the source. It helps when the source contributes more positively to the target margin than the displaced donor, unless downstream curvature overwhelms that gain. Source-max can therefore succeed, saturate, or fail under the same attention-mass rule. The supplementary material derives Equation 5, extends it to evidence sets, and states a sufficient margin certificate.

4 Experimental Setup

Controlled tasks. Argmax retrieval presents key-value pairs and requests the value paired with the largest key. Flag retrieval presents values with exactly one marked position and requests that value. Both have one labeled source. Two-evidence modular sum marks two digits and requests their sum modulo ten, so both sources are required. Four-layer decoders are trained with NoPE and RoPE and four seeds per cell; evaluation reaches $50\times$ training length. An eight-layer, width-512 grid checks scale. The inverse temperature $\beta \in \{1, 2, 4\}$ changes attention capacity while each assignment comparison preserves the resulting sorted spectrum.

Pretrained and natural-language tasks. In-context key-value recall lists pairs followed by a query key; the paired value is the labeled source. We test Qwen2.5, Pythia, Gemma, and SmoLM2 across increasing pair counts. The selector audit compares circuits ranked on disjoint calibration data by source mass or by output-conditioned utility. Natural MCQA uses four-choice questions with one gold passage and unrelated passages. Questions must be correct with the full context and gold passage alone, but incorrect without context. This filtering defines a model-specific context-dependent subset and is not intended to estimate average performance over the full dataset. Its length ladder freezes the four-passage circuit before forming nested 8-, 16-, and 32-passage contexts with additional unrelated passages.

Interventions and estimation. The distractor control swaps the largest and smallest non-source weights, preserving source mass while permuting the same attention row. For a source position s , let $d^+ = \arg \max_j a_j$, $d^- = \arg \min_j a_j$, and $D = \{j : j \neq s\}$ over valid tokens. Source-max swaps (s, d^+) and source-min swaps (s, d^-) . In controlled and pretrained key-value recall, the distractor control uses the maximum and minimum within D . Natural QA uses a displacement-matched control. With $\delta = a_{d^+} - a_s$, it chooses

$$(p, q) = \arg \min_{\{j, k\} \subset D} |a_j - a_k| - \delta \quad (6)$$

155 and swaps (p, q) , minimizing the L_1 -displacement mismatch to the treatment while holding source
 156 mass fixed; every control preserves the complete spectrum. The control does not match the donor to-
 157 ken, positional distance, or value norm, which remain separate confounds. Calibration and evaluation
 158 examples are disjoint, circuits are frozen before paired evaluation, hierarchical intervals resample
 159 calibration or model seeds above paired examples, and a cell is interpreted only after its frozen
 160 competence gate passes. Three confirmatory contrasts form one inferential family (controlled source-
 161 max versus distractor control, the capacity-by-assignment interaction, and natural utility-selected
 162 source-max versus matched control), tested with two-sided exact sign-flip tests at the cluster level
 163 with Holm correction; other intervals are descriptive and unadjusted. These swaps are counterfactual
 164 probes and need not lie on the model’s endogenous query-key manifold. Treatment and control share
 165 this intervention class, so their contrast measures downstream sensitivity to assignment, rather than
 166 the frequency with which the model naturally creates a particular row. Layer controls (+1.107 for se-
 167 lected $K = 8$ versus -0.059 for a random layer) and coverage dose response show greater sensitivity
 168 under this probe. They do not show that normal query-key computation produces swapped rows or
 169 that selected heads are the exclusive retrieval mechanism. Section 5.4 reports the endogenous-row
 170 test and the patching comparison that bound this reading.




Figure 2: Paired causal protocol. A circuit selected on disjoint calibration data is frozen before treatment and control are applied to the same evaluation row. Both paths preserve the spectrum; the natural QA control additionally minimizes displacement mismatch in Equation 6.

171 **Measured outcomes.** We record task accuracy, target-versus-competitor margin, source mass,
 172 attention-output variance, entropy, norms, and the full sorted spectrum. Every fixed-spectrum
 173 condition preserves the selected attention rows to numerical precision. The supplementary material
 174 reports prompts, hyperparameters, per-length tables, competence failures, and preregistered decision
 175 rules.




Figure 3: Assignment changes retrieval despite an identical attention spectrum. Source-max improves accuracy, source-min harms it, and the distractor-only control remains near zero (a). The effect increases with the fraction of the selected circuit receiving the intervention (b). Diamonds and intervals summarize trained models; faint paths show individual cells.




Figure 4: Capacity controls leverage, and utility predicts the finite intervention. Sharper spectra increase the correct-versus-wrong assignment contrast while sharpening alone remains near zero (a). The first-order utility term predicts the exact source-max margin effect across tasks and positional encodings (b). Error bars summarize trained models.

176 5 Results and Discussion

177 5.1 Assignment Changes Accuracy at Fixed Spectrum

178 We begin with the simplest test: move the largest available weight to the source, move it away, or
 179 exchange two distractors. Across sixteen trained models, the mean changes are +0.153, −0.043, and
 180 +0.011, respectively (Figure 3). Source-max beats the control in every model-length comparison.
 181 Their primary cluster-level contrast is +.142 (95% CI [+0.080, +.206]; exact $p = 3.05 \times 10^{-5}$;
 182 Holm-adjusted $p = 9.16 \times 10^{-5}$), while complete-spectrum invariant error stays below 10^{-6} .

183 The effect grows with circuit coverage (Figure 3b): patching one, two, four, and eight selected
 184 heads gives mean source-max gains of 0.007, 0.013, 0.056, and 0.153, monotone in 31 of 32
 185 comparisons, consistent with a distributed effect. Source attention tracks accuracy more closely than
 186 attention-output variance: normalizing that variance leaves accuracy unchanged, whereas forcing
 187 mass onto the known source restores controlled retrieval, and sharpening the original row has little
 188 effect. The ordering survives the eight-layer replication, both positional encodings, four seeds, and
 189 both controlled tasks, while random-layer and adjacent-layer interventions stay near zero. The full
 190 displacement audit, per-cell intervals, and variance and patching results are in the supplementary
 191 material.

192 5.2 Capacity Amplifies Assignment

193 We next vary spectrum and assignment independently. Temperature scaling changes the maximum
 194 available weight, and at every temperature we evaluate identity, source-max, source-min, and dis-
 195 tractor control on the same examples. The source-max minus source-min contrast grows from 0.196
 196 to 0.346, while sharpening the identity assignment changes accuracy by at most 0.009 (Figure 4a).
 197 Capacity sets the possible leverage, while assignment determines where that leverage lands. The




Figure 5: Matched pretrained-model comparison. A single calibration seed (seed 0) is shown because the frozen circuit is selected once per seed and the panels compare conditions within that fixed circuit; seed-pooled results are in the supplementary material. Error bars are paired 95% intervals over $n = 128$ examples per model-length cell. All panels use the 5-, 20-, and 80-pair ladder and report source-max, source-min, and untouched margin relative to the same distractor control. Qwen and Pythia respond positively, Gemma is saturated, and the source-mass-selected SmoLM2 circuit does not benefit.

198 primary interaction is $+0.150$ (95% CI $[+0.094, +0.210]$; exact $p = 6.10 \times 10^{-5}$; Holm-adjusted
199 $p = 1.22 \times 10^{-4}$).

200 For held-out examples we differentiate the target margin along the exact source-max path. The first-
201 order term in Proposition 2 ranks the finite margin change, with per-model Spearman correlations
202 of 0.87, 0.54, and 0.55 on Qwen2.5-1.5B, Pythia-1.4B, and Gemma-2-2B and sign agreements of
203 0.81, 0.81, and 0.71 (Appendix A). Aggregating to the 32 controlled cell means in Figure 4b raises
204 the rank correlation to 0.989, but that pooled figure overstates the per-row relation and we do not
205 use it as the headline. The slope of the actual-on-predicted fit is 1.387, off by about 39%, so we read
206 the utility term as a usable local ranking signal rather than a calibrated estimator of effect size. The
207 remaining error measures downstream curvature along the swap.

208 Modular sum requires the model to route and combine two marked values. At $5\times$ training length,
209 assigning the top two weights to evidence raises exact match by 0.066, assigning the bottom two
210 lowers it by 0.291, and evidence-max exceeds a control that preserves both the spectrum and total
211 evidence mass by 0.043 (95% CI $[0.017, 0.072]$). This is a small but positive two-evidence effect at
212 $5\times$ (the effect at $1\times$ is $+0.007$), so the result reaches two required tokens rather than one; it does
213 not establish general multi-token routing. A registered three-evidence estimate includes zero, and
214 four-evidence models miss the frozen competence gate, so arity beyond two remains unresolved.

215 Evidence-max and evidence-min preserve the complete spectrum. The multi-evidence control also
216 preserves total evidence mass, changing only which evidence token receives each weight. Its small
217 effect separates aggregate source mass from useful evidence assignment.

218 5.3 The Same Intervention Produces Three Pretrained Regimes

219 We read the pretrained results as directional, not as calibrated magnitudes, because the magnitudes
220 are seed-sensitive. The one condition we replicated across seeds is telling: on Qwen, the headline
221 seed selects a layer-19 circuit with a source-max margin of $+1.882$, while five further seeds all select
222 a layer-14 circuit with margins near $+0.04$ (mean $+0.042$), a $45\times$ gap. The six-seed mean is $+0.349$
223 with a hierarchical 95% interval of $[+0.037, +0.982]$ that excludes zero, so the direction is reproducible
224 while the size is not, and the single-seed magnitudes in Table 2 should be read as upper estimates.
225 The robust reading is the sign and the regime pattern: Qwen and Pythia gain margin from source-max,
226 Gemma stays accurate and gains little from extra source mass while source-min still harms it (the
227 saturation regime of Proposition 2), and the source-mass-selected SmoLM2 circuit does not benefit.

228 SmoLM2’s source-mass selector fails, yet ranking heads by an output-conditioned score restores a
229 positive effect. We are cautious about how much this isolates: our utility term is the gradient of the
230 very margin we then measure, so selecting heads by predicted ΔM and finding large ΔM is close to

Task	Cond.	Unpatched	Δ src-max	Δ src-min	Δ distr.	max—distr.
Single evid. (Acc)	$\beta = 1$	.619	+ .153	− .043	+ .011	+ .142
	$\beta = 4$	.619	+ .298	− .048	+ .025	+ .273
Two evid. (EM)	$1 \times$	.952	+ .019	− .559	+ .012	+ .007
	$5 \times$	.484	+ .066	− .291	+ .023	+ .043

Table 1: Controlled fixed-spectrum assignment. Single-evidence rows report per-token accuracy over $n = 8,192$ head-example pairs; two-evidence rows report exact match over $n = 2,048$. Source-max moves the row maximum to the source, source-min moves it away, and the distractor control swaps two non-source weights; all preserve the complete spectrum. Full cells and 95% intervals are in the supplementary material.

Model	Unpatched	Δ src-max	Δ src-min	max—distr.
Qwen2.5-1.5B [†]	−3.048	+1.990	− .961	+1.799
Pythia-1.4B	−3.805	+ .826	−1.022	+ .680
Gemma-2-2B	+2.234	+ .040	−1.373	+ .027
SmolLM2-1.7B	−2.782	− .168	− .112	− .091
<i>Output-conditioned circuit selection (ΔM vs. matched control)</i>				
	util. gain	util. gap	src mass	util—src
Pythia-1.4B	+ .519	+ .359	+ .006	+ .513
SmolLM2-1.7B	+1.253	+ .859	+ .006	+1.247
Qwen2.5-1.5B	+1.493	+1.493	+ .069	+1.424

Table 2: Pretrained models, target-margin units ($n = 384$ per model, matched 5/20/80-pair ladder from Figure 5). **These are single-calibration-seed point estimates and should be read as directional, not calibrated:** the only condition we replicated across seeds ([†], Qwen top-block source-max) shows the headline seed is $45\times$ the five-seed mean (six-seed mean $+0.349$, 95% CI $[+.037, +.982]$; Appendix A), so the other single-seed magnitudes are likely inflated. Top: fixed-spectrum assignment. Bottom: output-conditioned circuit selection; because the utility term is the gradient of this margin, its advantage over source mass (src mass column) shows that outcome-aware beats outcome-blind selection rather than isolating utility. A Gemma selector run was not conducted. Full cells and 95% intervals are in the supplementary material.

231 guaranteed by Proposition 2, and disjoint calibration data controls overfitting but not this circularity
232 of functional form. The supported claim is therefore the weaker one that *outcome-aware* selection
233 beats outcome-blind source-mass selection; isolating a special role for the utility functional would
234 require an outcome-aware but non-utility baseline, such as attribution-patching or ablation-based head
235 scores, which we have not run. On context-grounded natural QA the same pattern holds directionally
236 across six seeds (mean contrast $+ .378$, one seed dominant, leave-one-seed-out $+ .19$ to $+ .55$); we
237 report the interval rather than a significance claim, since an exact sign test at six seeds bottoms out at
238 the $p = .03125$ we reach.

239 5.4 Two results that bound the causal claim

240 Two measurements limit how far the fixed-spectrum effect should be read as the model’s own
241 mechanism, and we state them here rather than as asides.

242 First, the effect is measured on rows the model does not itself produce. On untouched distractor-order
243 variants, where we do not swap, the conditional source-mass coefficient is $+1.250$ ($p = .03125$),
244 so source mass is associated with the outcome in endogenous data. The matched fixed-spectrum
245 contrast on those same endogenous rows, however, is $+ .026$ with a 95% interval of $[-.017, +.070]$
246 that includes zero. The intervention thus measures sensitivity to a counterfactual reassignment, and it
247 does not establish that normal query-key computation creates the swapped rows. This is the central
248 caveat of an off-manifold probe, and the near-zero endogenous contrast is why we frame FSS as a
249 test of a specified reassignment rather than of the model’s spontaneous routing.

250 Second, restoring the selected circuit recovers little of the damage from corrupting it. Clean-to-corrupt
251 activation patching on the selected heads rescues $+ .308$ margin, which is only 2.5 to 3.3% of the




Figure 6: Matched natural-QA length ladder (seed 0; the four-passage circuit is frozen once per seed and reused at every length, and the additional Qwen seed is in the supplementary material). Bands are paired bootstrap 95% intervals over $n = 128$ examples per cell. From 4 to 32 passages, accuracy falls (a), source routing moves in opposite directions across models (b), and frozen-circuit rescue stays flat or weakens (c).

median corruption damage. Because patching restores head outputs after source-value corruption while reassignment acts on an unchanged input, the two magnitudes are not directly comparable; still, the small recovered fraction indicates that the selected heads are not where most of the causal action sits, and the effect we report is one contribution among several rather than the dominant pathway.

5.5 Natural Context Marks the Scope Boundary

At a fixed context in ordinary text, the answer is sensitive to a spectrum-preserving reassignment onto the gold evidence. We stop short of “causal routing to the required evidence”: without the semantic-distractor control (Limitations), the result is still compatible with sensitivity to any answer-shaped token, and that control is the precondition for the stronger claim. The length ladder is a scope test: we select the circuit at four passages, freeze every choice, and append unrelated passages to form nested 4-, 8-, 16-, and 32-passage contexts. Both Qwen and SmoLM2 lose margin and accuracy by 32 passages, yet their source-mass trajectories diverge: Qwen source mass rises as rescue weakens, while SmoLM2 source mass falls with rescue flat. The shared degradation cannot be attributed to one consistent source-mass trajectory, so this is a scope limit of the diagnosis, not a failed prediction. At a fixed context the swap measures whether a chosen circuit can use reassigned evidence; as contexts grow, representation quality, distractor semantics, and computation outside the circuit also matter, and the theorem, which concerns a specified swap at a fixed context, does not resolve those bottlenecks. Exact 32-minus-4 changes are in the supplementary material.

5.6 Identification and Replication Audits

Replication checks vary the calibration sample, patch budget, and model family, and they support the same reading. Selected circuits are positive at every head budget, adjacent and random layers are smaller or near zero, the utility term ranks exact effects within every competent cell while calibrating them imperfectly, and an exact rerun reproduces the direction with an atypically large original magnitude. Appendix A reports these audits in full.

6 Limitations

FSS identifies the causal effect of a specified post-softmax reassignment, not a complete account of model computation: the swap is an off-manifold counterfactual, and preserving the spectrum isolates assignment without showing that training produces the same row. One control is also missing and bounds the natural-QA reading. The displacement-matched control moves mass between two task-irrelevant tokens, whereas the treatment moves it onto a semantically loaded token; a stronger control moves the same maximum onto a competitor answer token, matched for spectrum and displacement, to separate assignment to required evidence from assignment to any answer-shaped token (source-min is not a substitute, since it also changes the magnitude moved). We flag this semantic-distractor control as the primary missing comparison, alongside adding calibration seeds so the pretrained cells report distributions, not point estimates. Higher arity is competence-limited and the generation audit is short-horizon. Anonymized code, seeds, prompts, and per-example outputs accompany the submission and will be released on acceptance.

References

- Gino Brunner, Yang Liu, Damián Pascual, Oliver Richter, Massimiliano Ciaramita, and Roger Wattenhofer. On identifiability in transformers. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=BJg1f6EFDB>.
- Arthur Conmy, Augustine Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/34e1dbe95d34d7ebaf99b9bcae5b2be-Abstract-Conference.html.
- Harry Jake Cunningham and Nicola Muca Cirone. Contribution weights: A geometrical analysis of self-attention transformers. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026. URL <https://arxiv.org/abs/2606.07604>.
- Mor Geva, Jasmijn Bastings, Katja Filippova, and Amir Globerson. Dissecting recall of factual associations in auto-regressive language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12216–12235. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.emnlp-main.751. URL <https://aclanthology.org/2023.emnlp-main.751/>.
- G. H. Hardy, J. E. Littlewood, and G. Pólya. *Inequalities*. Cambridge University Press, 2 edition, 1952.
- Michael Hassid, Hao Peng, Daniel Rotem, Jungo Kasai, Ivan Montero, Noah A. Smith, and Roy Schwartz. How much does attention actually attend? questioning the importance of attention in pretrained transformers. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, 2022. URL <https://aclanthology.org/2022.findings-emnlp.101/>.
- Xiang Hu, Zhihao Teng, Jun Zhao, Wei Wu, and Kewei Tu. Efficient length-generalizable attention via causal retrieval for long-context language modeling. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 24727–24743. PMLR, 2025. URL <https://proceedings.mlr.press/v267/hu25v.html>.
- Sarthak Jain and Byron C. Wallace. Attention is not explanation. In *Proceedings of NAACL-HLT*, pages 3543–3556, 2019. URL <https://aclanthology.org/N19-1357/>.
- Amirhossein Kazemnejad, Inkit Padhi, Karthikeyan Natesan Ramamurthy, Payel Das, and Siva Reddy. The impact of positional encoding on length generalization in transformers, 2023.
- Goro Kobayashi, Tatsuki Kuribayashi, Sho Yokoi, and Kentaro Inui. Attention Is Not Only a Weight: Analyzing transformers with vector norms. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 7057–7075. Association for Computational Linguistics, 2020. doi: 10.18653/v1/2020.emnlp-main.574. URL <https://aclanthology.org/2020.emnlp-main.574/>.
- János Kramár, Tom Lieberum, Rohin Shah, and Neel Nanda. AtP*: An efficient and scalable method for localizing LLM behaviour to components, 2024. URL <https://arxiv.org/abs/2403.00745>.
- Ruining Li, Gabor Boduljak, and Jinghao Zhou. On vanishing variance in transformer length generalization, 2025.
- Omar Naim and Nicholas Asher. On explaining with attention matrices. In *ECAI 2024*, pages 1035–1042. IOS Press, 2024. doi: 10.3233/FAIA240594. URL <https://doi.org/10.3233/FAIA240594>.
- Ken M. Nakanishi. Scalable-softmax is superior for attention, 2025.
- Yurii Nesterov. *Introductory Lectures on Convex Optimization: A Basic Course*. Kluwer Academic Publishers, 2004.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392, 2016. doi: 10.18653/v1/D16-1264.
- Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Thomas Adler, Lukas Gruber, Markus Holzleitner, Milena Pavlović, Geir Kjetil Sandve, Victor Greiff, David Kreil, Michael Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. Hopfield networks is all you need. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=tL89RnzIiCd>.
- Aaquib Syed, Can Rager, and Arthur Conmy. Attribution patching outperforms automated circuit discovery, 2023. URL <https://arxiv.org/abs/2310.10348>.

- 340 Davoud Ataee Tarzanagh, Yingcong Li, Xuechen Zhang, and Samet Oymak. Max-margin to-
 341 ken selection in attention mechanism. In *Advances in Neural Information Processing Systems*,
 342 volume 36, 2023. URL [https://proceedings.neurips.cc/paper_files/paper/2023/hash/](https://proceedings.neurips.cc/paper_files/paper/2023/hash/970f59b22f4c72aec75174aae63c7459-Abstract-Conference.html)
 343 [970f59b22f4c72aec75174aae63c7459-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/970f59b22f4c72aec75174aae63c7459-Abstract-Conference.html).
- 344 Pavlo Vasylenko, Hugo Pitorro, André F. T. Martins, and Marcos Treviso. Long-context generalization with
 345 sparse attention, 2025.
- 346 Jie Wang, Tao Ji, Yuanbin Wu, Hang Yan, Tao Gui, Qi Zhang, Xuanjing Huang, and Xiaoling Wang. Length
 347 generalization of causal transformers without position encoding, 2024.
- 348 Sarah Wiegrefe and Yuval Pinter. Attention is not not explanation. In *Proceedings of EMNLP-IJCNLP*, pages
 349 11–20, 2019. URL <https://aclanthology.org/D19-1002/>.
- 350 Takateru Yamakoshi, James L. McClelland, Adele E. Goldberg, and Robert D. Hawkins. Causal interventions
 351 expose implicit situation models for commonsense language understanding. In *Findings of the Association for*
 352 *Computational Linguistics: ACL 2023*, pages 13265–13293, 2023. doi: 10.18653/v1/2023.findings-acl.839.
 353 URL <https://aclanthology.org/2023.findings-acl.839/>.
- 354 Pál Zsámboki, Benjamin Levi, David Ansel Josef Smith, Mitansh Kagalwala, Arlington Kell, Samuel Liechty,
 355 and Cong Wang. Post-norm can resharpen attention, 2025.

356 A Identification and replication audits

357 Replication checks vary the calibration sample, patch budget, and model family. Across three frozen
 358 Qwen calibration seeds, selected circuits are positive at every head budget: $K = 2, 4, 8$ gives $+0.460$,
 359 $+1.054$, and $+1.107$ margin. Adjacent layers are smaller and inconsistent ($+0.147$ and -0.131), while
 360 random layers average near zero (-0.059). Full intervals and seed standard deviations are reported in
 361 the supplementary package.

362 A separate replication tests the utility relation. Circuits are chosen by source mass on disjoint
 363 calibration examples; utility is then computed on held-out evaluation rows. Across three seeds per
 364 family, every competent source-max cell is positive, and predicted effects rank the exact effects
 365 within each cell. Mean Spearman correlations are $.866$ for Qwen2.5-1.5B, $.544$ for Pythia-1.4B, and
 366 $.554$ for Gemma-2-2B, with sign agreements $.805$, $.807$, and $.714$. The approximation is imperfect at
 367 large effects, so we use it as a local ranking signal rather than a calibrated finite-effect estimator. For
 368 the 32 controlled cell means in Figure 4b, the actual-on-predicted fit is $1.387x - 0.102$ (Pearson $.978$),
 369 which deviates from identity calibration.

370 The exact Qwen rerun, with precision, lengths, contrast, and sample counts fixed, is the sharpest
 371 evidence that the pretrained magnitudes are seed-sensitive and the main reason we present them
 372 as directional. Each seed selects its own circuit before evaluation, and the headline seed selects a
 373 different circuit from the other five (Table 3). The sign is reproducible in all six seeds, the six-seed
 374 mean source-max minus control margin is $+0.349$ with a hierarchical 95% interval of $[+0.037, +0.982]$
 375 that excludes zero, but the size varies by more than an order of magnitude. A two-sided exact sign test
 376 at six seeds bottoms out at $p = .03125$, so we report the seed-level values rather than a significance
 377 verdict. Adding seeds to every pretrained cell, so that Table 2 reports distributions rather than point
 378 estimates, is the most important open item, not a cosmetic extension.

Seed	Layer	Heads	Mean margin
0 (headline)	19	[3, 1, 0, 2]	+1.882
1	14	[3, 4, 0, 1]	+0.041
2	14	[3, 4, 0, 1]	+0.036
3	14	[3, 0, 4, 1]	+0.053
4	14	[3, 0, 4, 1]	+0.039
5	14	[3, 0, 4, 1]	+0.044

Table 3: Exact multi-seed Qwen replication. The headline seed selects a layer-19 circuit and a margin $45\times$ the mean of the five replication seeds, which agree on a layer-14 circuit near $+0.04$. The effect is directionally reproducible and magnitude-unstable.

Output-conditioned selectors lead source-mass selection in each family, but this comparison is outcome-aware versus outcome-blind and does not isolate the utility functional, which is the gradient of the measured margin; an outcome-aware but non-utility baseline is not run here. Full selector results, competence-gate failures, and the anonymized code package appear in the supplementary material.

B Theory details

This appendix gives the full statements and proofs behind Propositions 1 and 2, and a multi-evidence corollary. The complete per-cell result tables and secondary figures are in the separate supplementary material.

Established softmax grouping identity. For evidence positions S , distractors D , scores s_j , and temperature $T > 0$, let $Z_G = \sum_{j \in G} e^{s_j/T}$ and $F_G = -T \log Z_G$. Grouping the softmax denominator gives

$$m_S := \sum_{j \in S} a_j = \frac{Z_S}{Z_S + Z_D} = \sigma\left(\frac{F_D - F_S}{T}\right). \quad (7)$$

With the group log-mean-exp score $\bar{s}_G = T \log(|G|^{-1} \sum_{j \in G} e^{s_j/T})$,

$$m_S \geq \alpha \iff \bar{s}_S - \bar{s}_D \geq T \left(\log \frac{|D|}{|S|} + \text{logit}(\alpha) \right). \quad (8)$$

These Gibbs partition-function identities [Ramsauer et al., 2021] make the dependence on distractor scores explicit: distractor count alone is insufficient, since arbitrarily many negligible-score distractors need not reduce evidence mass.

Capacity-assignment decomposition. Let $a_{(1)} \geq \dots \geq a_{(n)}$, let S contain r evidence positions, and define $C_r^+ = \sum_{i=1}^r a_{(i)}$, $C_r^- = \sum_{i=n-r+1}^n a_{(i)}$, and $K_r(a) = C_r^+ - C_r^-$.

Theorem 1 (Capacity-assignment decomposition). *For every permutation π of a fixed attention-weight multiset, $C_r^- \leq m_S(\pi) \leq C_r^+$, both endpoints are attainable, and for $K_r(a) > 0$ the unique coordinate $\rho_S(\pi) = (m_S(\pi) - C_r^-)/K_r(a) \in [0, 1]$ gives $m_S(\pi) = C_r^- + \rho_S(\pi)K_r(a)$. For affine margin $M(\pi) = b + \sum_j a_{\pi(j)}u_j$ with sorted utilities $u_{(1)} \geq \dots \geq u_{(n)}$,*

$$\max_{\pi} M(\pi) = b + \sum_i a_{(i)}u_{(i)}, \quad \min_{\pi} M(\pi) = b + \sum_i a_{(i)}u_{(n+1-i)}. \quad (9)$$

No statistic invariant to token permutation can determine evidence assignment or the output margin.

Proof. An r -position set receives at most the r largest weights and at least the r smallest; assigning those weights to S attains the endpoints. Normalizing the bounded mass gives ρ_S . For the output range, suppose $u_p > u_q$ but $a_{\pi(p)} < a_{\pi(q)}$. Swapping those weights changes the margin by $(a_{\pi(q)} - a_{\pi(p)})(u_p - u_q) > 0$. Repeated exchanges align weights and utilities for the maximum and reverse them for the minimum, the classical rearrangement inequality [Hardy et al., 1952]. The extremal assignments preserve the complete weight multiset and therefore every permutation-invariant statistic. $\square$

Nonlinear downstream network. Fix an input and selected heads H at one layer. Let $x_{hj} = O_h v_{hj}$ and $z = \sum_{h \in H} \sum_j a_{hj} x_{hj}$. For target y and competitor c , let $M_c(z) = g_y(z) - g_c(z)$. A source-donor swap has $\delta_h = a_{hd_h} - a_{hs}$ and $\Delta z = \sum_{h \in H} \delta_h (x_{hs} - x_{hd_h})$. Define local utility $u_{hj}^{(c)} = \nabla M_c(z)^\top x_{hj}$.

Theorem 2 (Nonlinear routing-margin bound). *If ∇M_c is L_c -Lipschitz from z to $z + \Delta z$, then*

$$\left| \Delta M_c - \sum_{h \in H} \delta_h (u_{hs}^{(c)} - u_{hd_h}^{(c)}) \right| \leq \frac{L_c}{2} \|\Delta z\|_2^2. \quad (10)$$

415 *Proof.* The swap changes head h by exactly $\delta_h(x_{hs} - x_{hd_h})$. The fundamental theorem of calculus
 416 gives $\Delta M_c = \int_0^1 \nabla M_c(z + t\Delta z)^\top \Delta z dt$. Subtracting $\nabla M_c(z)^\top \Delta z$ and applying Lipschitzness
 417 bounds the remainder by $\int_0^1 L_c t \|\Delta z\|_2^2 dt = L_c \|\Delta z\|_2^2 / 2$ [Nesterov, 2004]. Expanding the first-order
 418 term gives Equation 10. $\square$

419 **Corollary 3** (Multi-evidence routing). *Let S contain r evidence positions and pair every evidence*
 420 *position missing a top- r weight with one top- r non-evidence donor. For disjoint swap pairs P_h and*
 421 *$\delta_{hsd} = a_{hd} - a_{hs} \geq 0$, $\Delta z = \sum_{h \in H} \sum_{(s,d) \in P_h} \delta_{hsd}(x_{hs} - x_{hd})$, and Equation 10 holds with the*
 422 *first-order sum extended over $(s, d) \in P_h$. Assigning top- r capacity to the evidence improves the*
 423 *margin whenever the summed evidence–donor utility gain exceeds the curvature remainder.*

424 The decomposition and corollary are algebraic for arbitrary r . Their behavioral premise is established
 425 for $r = 1$ and $r = 2$; the three-evidence interval includes zero and the four-evidence models miss the
 426 competence gate, so efficacy beyond $r = 2$ remains open. The condition is sufficient, not necessary:
 427 source-max can fail when evidence values have no greater downstream utility, patched heads cancel,
 428 or curvature dominates, so natural source mass alone is not a universal circuit-selection rule.

429 C Experimental protocol

430 **Controlled tasks and training.** Argmax retrieval requests the value paired with the largest key;
 431 flag retrieval requests the single marked value; both are order-invariant with one identifiable evidence
 432 position. Reversed decimal addition is an order-dependent contrast, and two-evidence modular
 433 sum marks two digits and requests their sum modulo ten so copying either digit is insufficient.
 434 The primary model is a four-layer pre-norm decoder transformer with width 256, eight heads, and
 435 feed-forward width 1024, trained with AdamW, batch size 512, and warmup plus cosine decay on
 436 lengths one through five. Evaluation follows a geometric ladder to $50\times$ training length, and the
 437 scale replication uses eight layers and width 512. A cell is interpreted only after training-length
 438 exact match reaches 0.8. Order-invariant cells use four seeds; the order-dependent contrast uses two.
 439 The preregistration fixes the competence gate, break-point rule, and variance-collapse check before
 440 intervention outcomes.

441 **Measured quantities and interventions.** At the answer query we record evidence mass, attention
 442 entropy, every attention norm, participation ratio, attention-output variance, and target margins.
 443 Source-max swaps the source weight with the row maximum, source-min swaps it with the row
 444 minimum, and the distractor control swaps the largest and smallest non-source weights. All conditions
 445 preserve model parameters, prompts, transformed values, the complete sorted attention row, entropy,
 446 norms, maximum, and participation ratio; the control additionally preserves source mass.

447 **Confirmatory inference.** One family of three primary contrasts is fixed in advance: controlled
 448 source-max minus distractor control, the capacity-by-assignment interaction, and natural utility-
 449 selected source-max minus matched control. Exact two-sided sign-flip tests operate on independent
 450 trained-model clusters for the controlled claims and calibration-seed clusters for natural QA, and
 451 Holm adjustment controls familywise error across the three tests. Hierarchical intervals and all other
 452 reported cells are descriptive and unadjusted unless stated otherwise.

453 D Scope, boundaries, and reproducibility

454 **Scope.** The equals-sign second-prompt format fails its competence gate in Qwen, and a blinded
 455 arrow-format pilot fails full-run competence despite directional margin signs; both are excluded
 456 from causal replications. Pythia and Qwen are floor-limited at their longest contexts, so margin
 457 is the interpretable tail outcome; Qwen-7B is quantized, and Gemma supplies a high-competence
 458 saturation boundary. The fixed-spectrum theorem still applies to the natural-QA length result because
 459 it conditions effects on transferred mass, value utility, and curvature, but the length ladder does not
 460 support the stronger claim that natural length degradation must involve source-mass loss.

461 **Reproducibility.** The controlled tasks are generated from the included task generator and require
 462 no external dataset. Natural MCQA samples are constructed from the public SQuAD training split

463 [Rajpurkar et al., 2016] with the sampling seed and eligibility rule included. Pretrained probes use
464 publicly released Qwen, Pythia, Gemma, and SmolLM2 checkpoints, with exact identifiers, tokenizer
465 settings, and precision recorded in the result artifacts. Pretrained experiments used a Google Colab
466 runtime with one NVIDIA T4 GPU, PyTorch, and Hugging Face Transformers; the earlier controlled
467 artifacts retain configurations and seeds but not the exact package-version snapshot, so we mark
468 infrastructure reproducibility partial. Saved artifacts reproduce every reported statistic and figure
469 without loading a language model, and an anonymized code and artifact archive accompanies the
470 submission as supplementary material.